COLUMBUS DAY 4.0

BIG DATA FOR .NET AND SQL

(AZURE DATA LAKE)

PLAN

- BIG DATA
- Azure Cloud
- Azure Data Lake
  - Azure Data Lake Store
  - Azure Data Lake Analytics
- U-SQL
- Azure Data Lake Live Azure Demo

BIG DATA

- 3V
- 4V (Veracity)

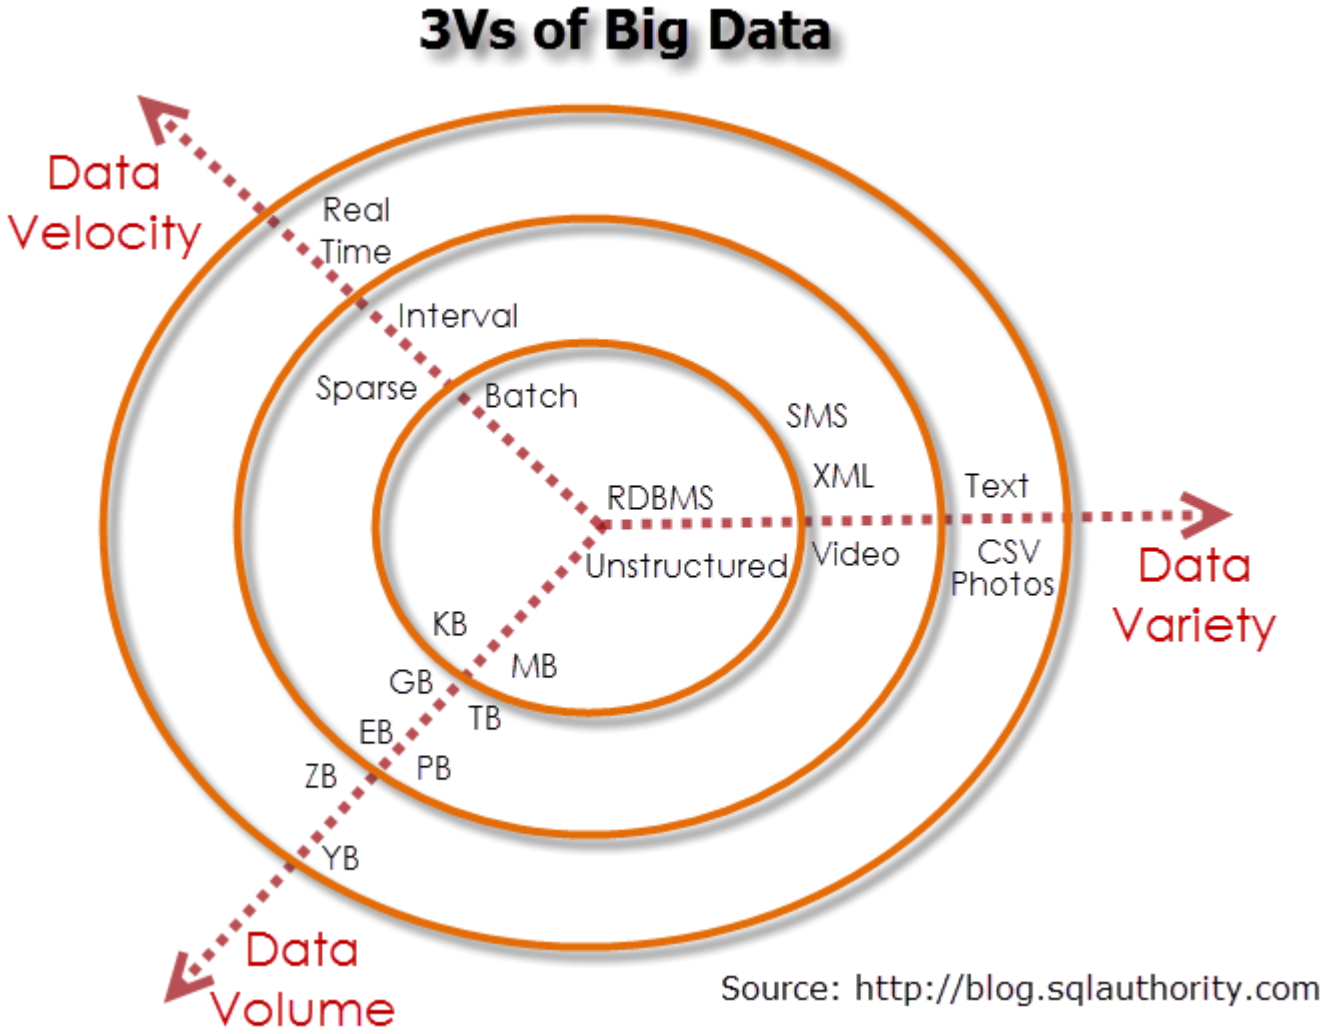

The Structure of Big Data

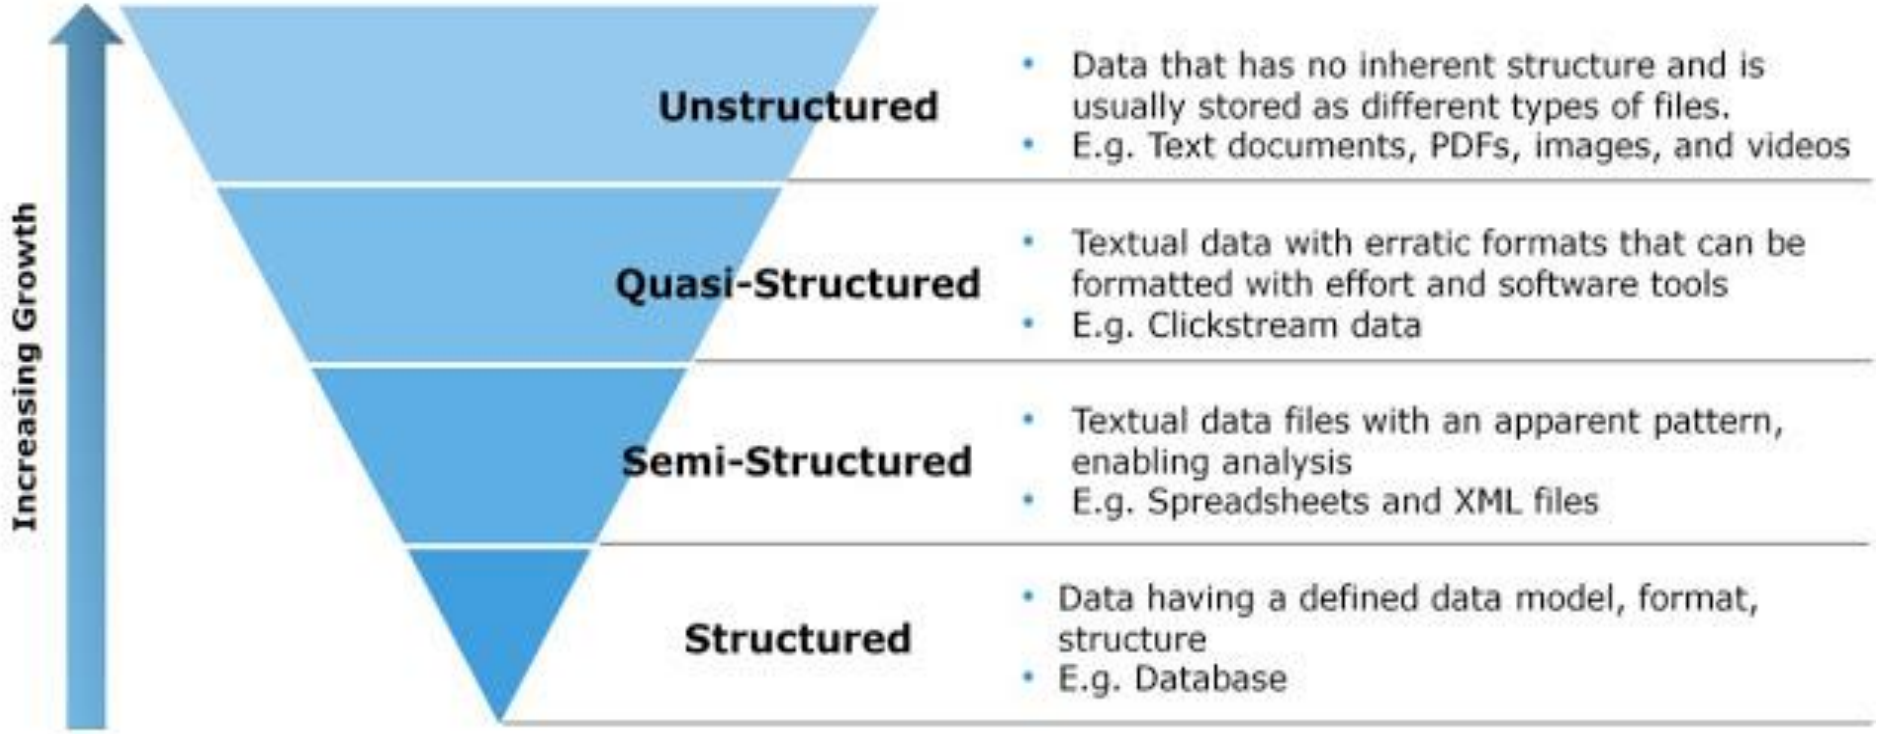

Source: http://www.tsmtutorials.com/2016/06/data-and-information-basics.html

Schema-on-Read vs Schema-on-Write

SCHEMA-ON-WRITE (RDBMS):

- Create static DB schema
- Transform data into RDBMS
- Query data in RDBMS format

New columns must be added explicitly before new data can propagate into the system.

SCHEMA-ON-READ (HADOOP OR ADLS):

- Copy data in its native format
- Create schema + parser
- Query Data in its native format (does ETL on the fly)

New data can start flowing any time and will appear retroactively once the schema/parser properly describes it.

Why Big Data?

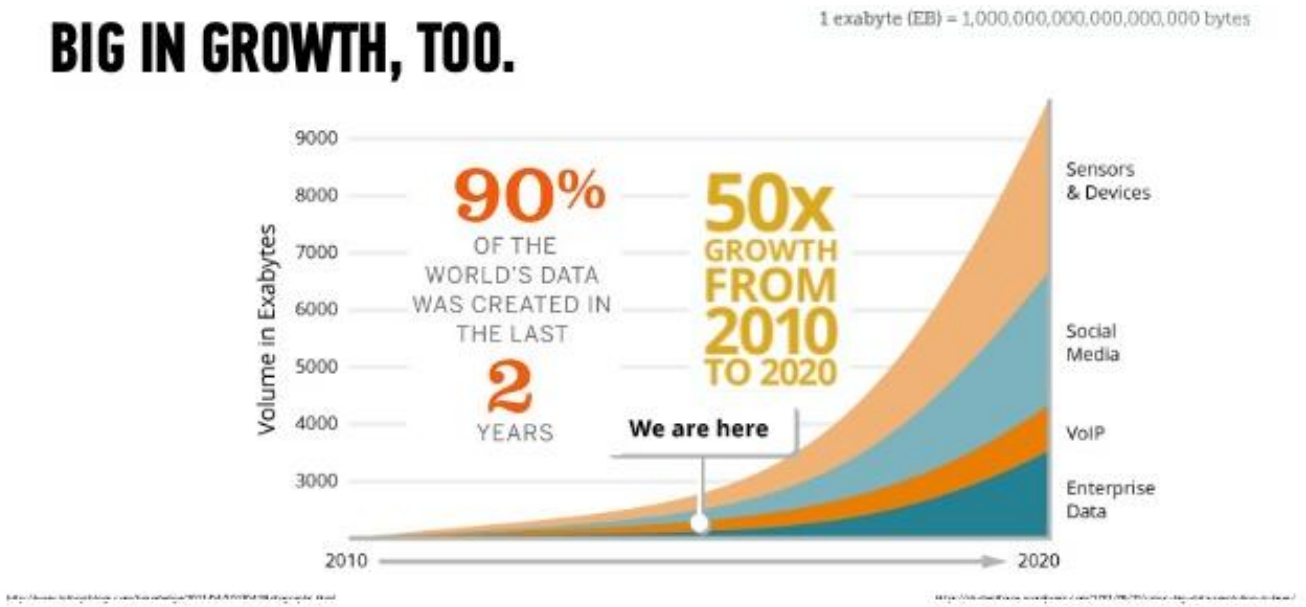
BIG IN GROWTH, TOO.
1 exabyte (EB) = 1,000,000,000,000,000,000 bytes
Volume in Exabytes
9000
8000
7000
6000
5000
4000
3000
90%
OF THE WORLD'S DATA WAS CREATED IN THE LAST
2
YEARS
50x
GROWTH FROM 2010 TO 2020
We are here
Sensors & Devices
Social Media
VoIP
Enterprise Data
2010
2020

Why Big Data?

	Rank	Brand	Brand Value	1-Yr Value Change	Brand Revenue	Company Advertising	Industry
	#1	Apple	$154.1 B	6%	$233.7 B	$1.8 B	Technology
	#2	Google	$82.5 B	26%	$68.5 B	$3.2 B	Technology
	#3	Microsoft	$75.2 B	9%	$87.6 B	$1.9 B	Technology
	#4	Coca-Cola	$58.5 B	4%	$21.9 B	$4 B	Beverages
	#5	Facebook	$52.6 B	44%	$17.4 B	$281 M	Technology

Source: http://www.forbes.com/powerful-brands/list/

Most Valuable Companies in the Fortune 500

MARKET VALUE RANK ▼	COMPANY	INDUSTRY	MARKET VALUE ($BIL)
1	Apple	Computers, Office Equipment	534
2	Alphabet	Internet Services and Retailing	507
3	Microsoft	Computer Software	413
4	Exxon Mobil	Petroleum Refining	326
5	Facebook	Internet Services and Retailing	321
6	Berkshire Hathaway	Insurance: Property and Casualty (Stock)	312
7	Johnson & Johnson	Pharmaceuticals	288
8	General Electric	Diversified Financials	271
9	Amazon.com	Internet Services and Retailing	250
10	Wells Fargo	Commercial Banks	242

Source: S&P Capital IQ

FORTUNE

Cloud Service Models and Azure

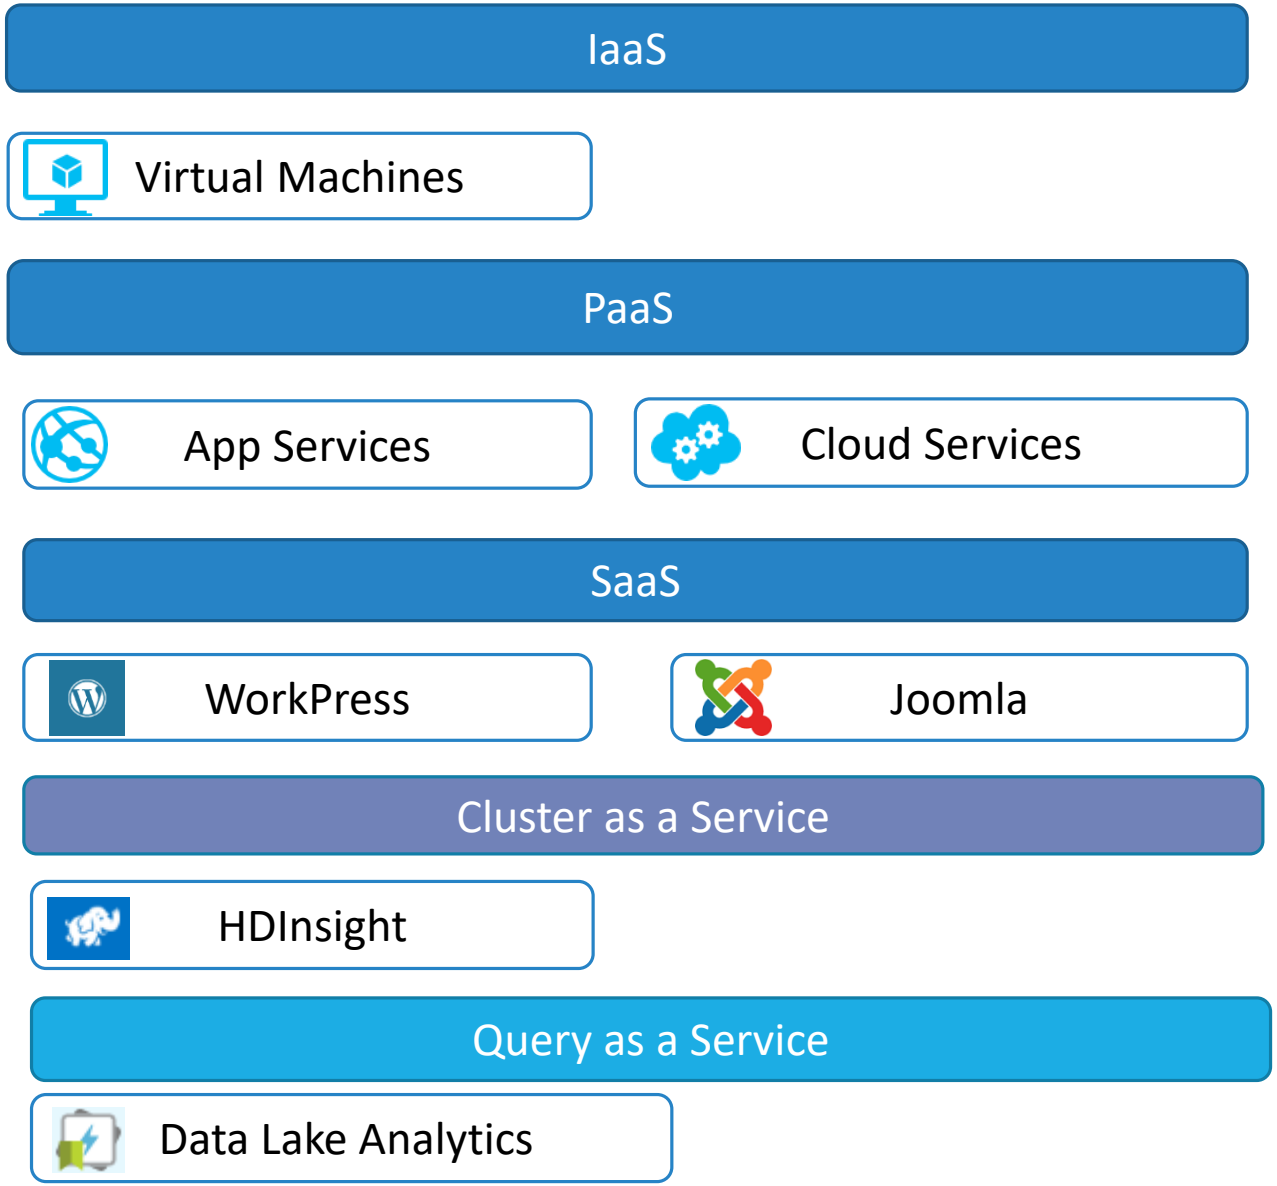

Azure Services

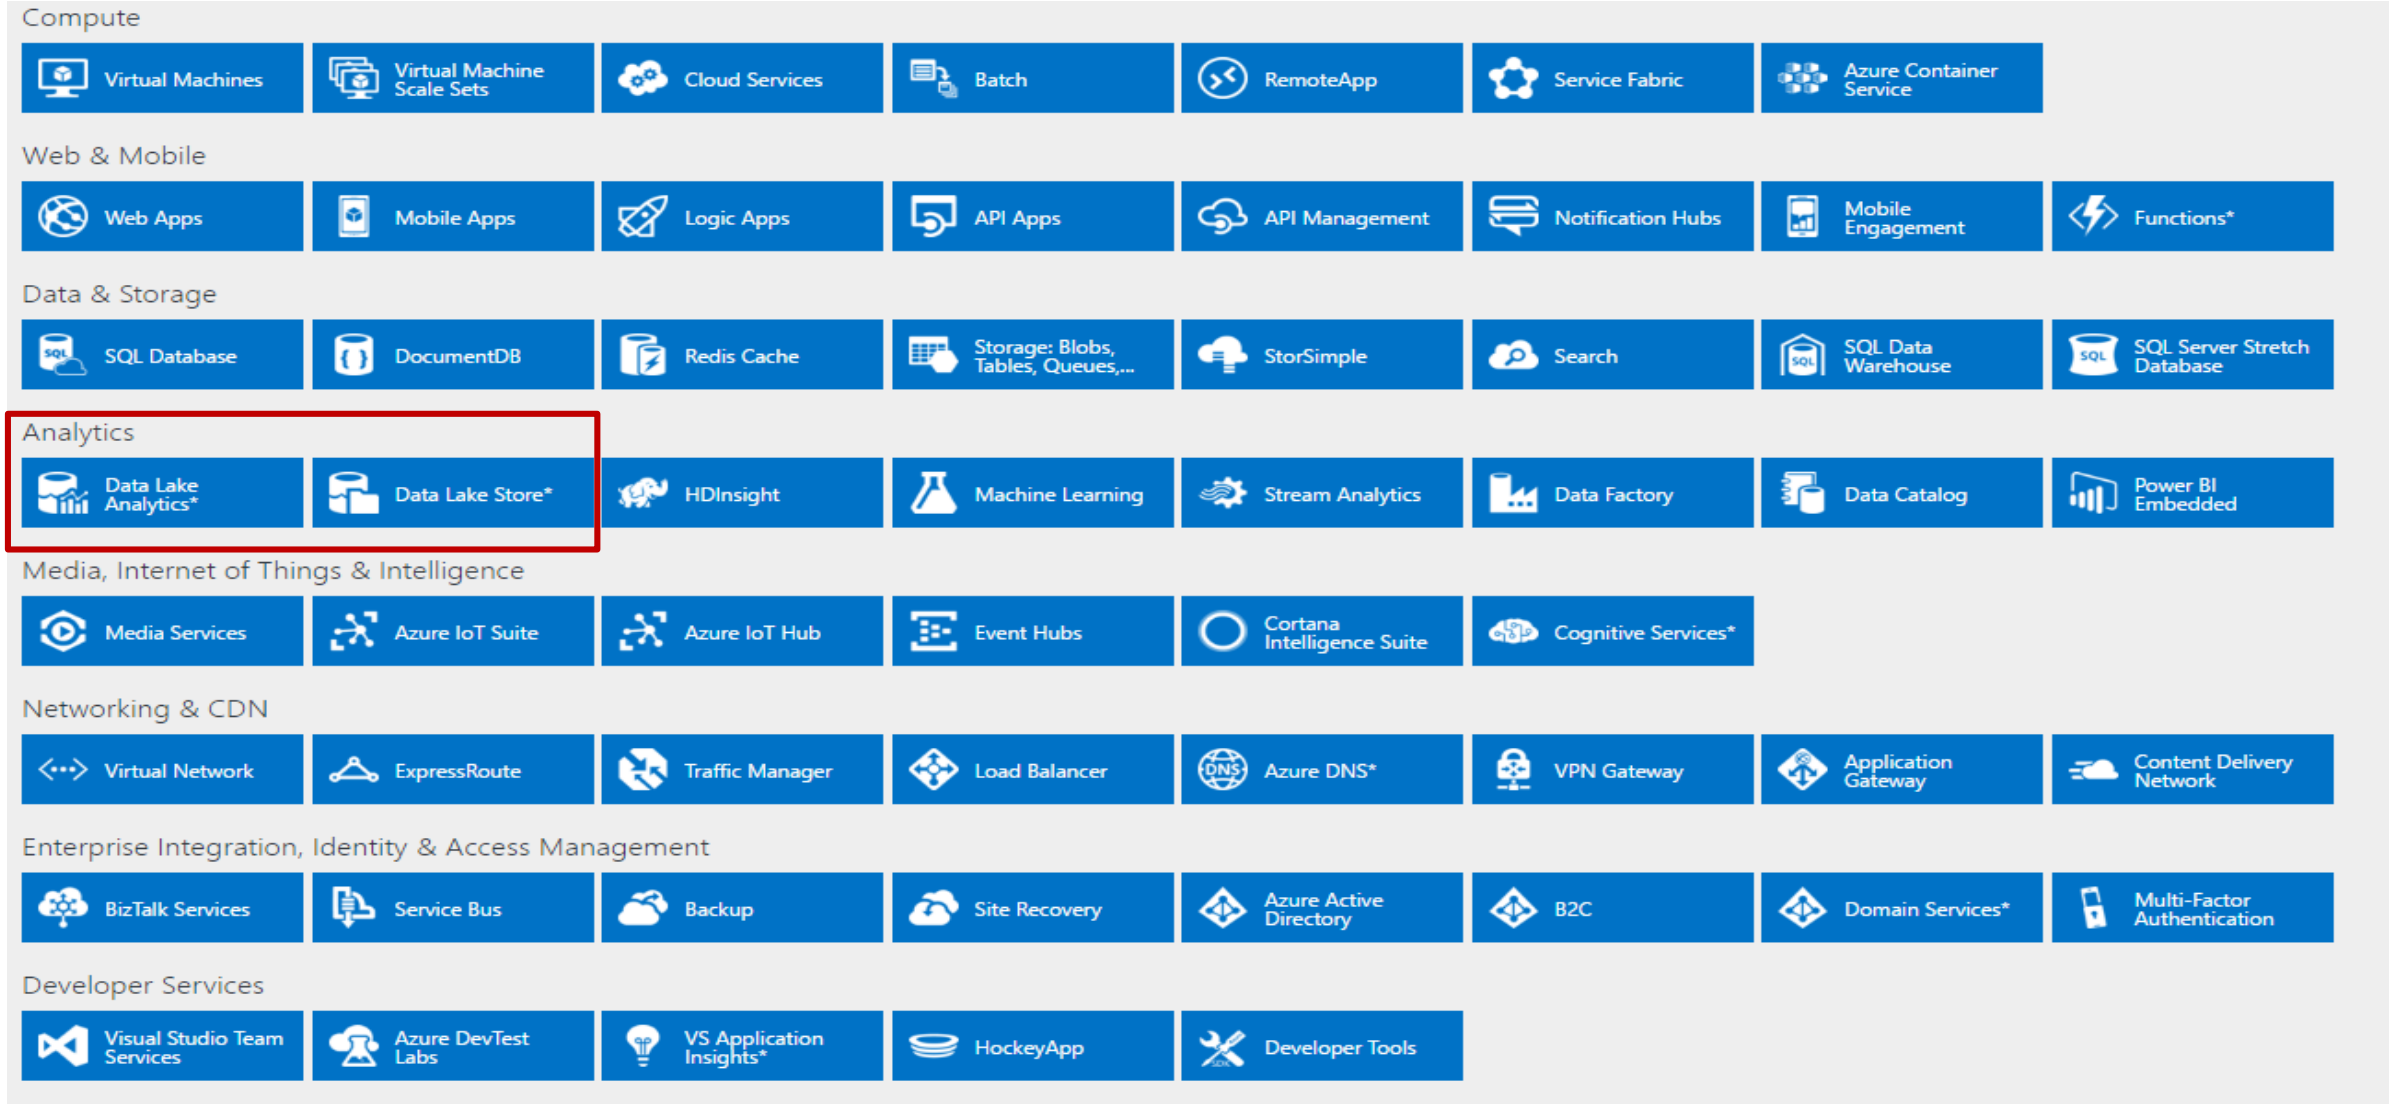

http://azureplatform.azurewebsites.net/en-us/

Azure Data Lake

Mike Rys

Saveen Reddy

Data Lake Appoach

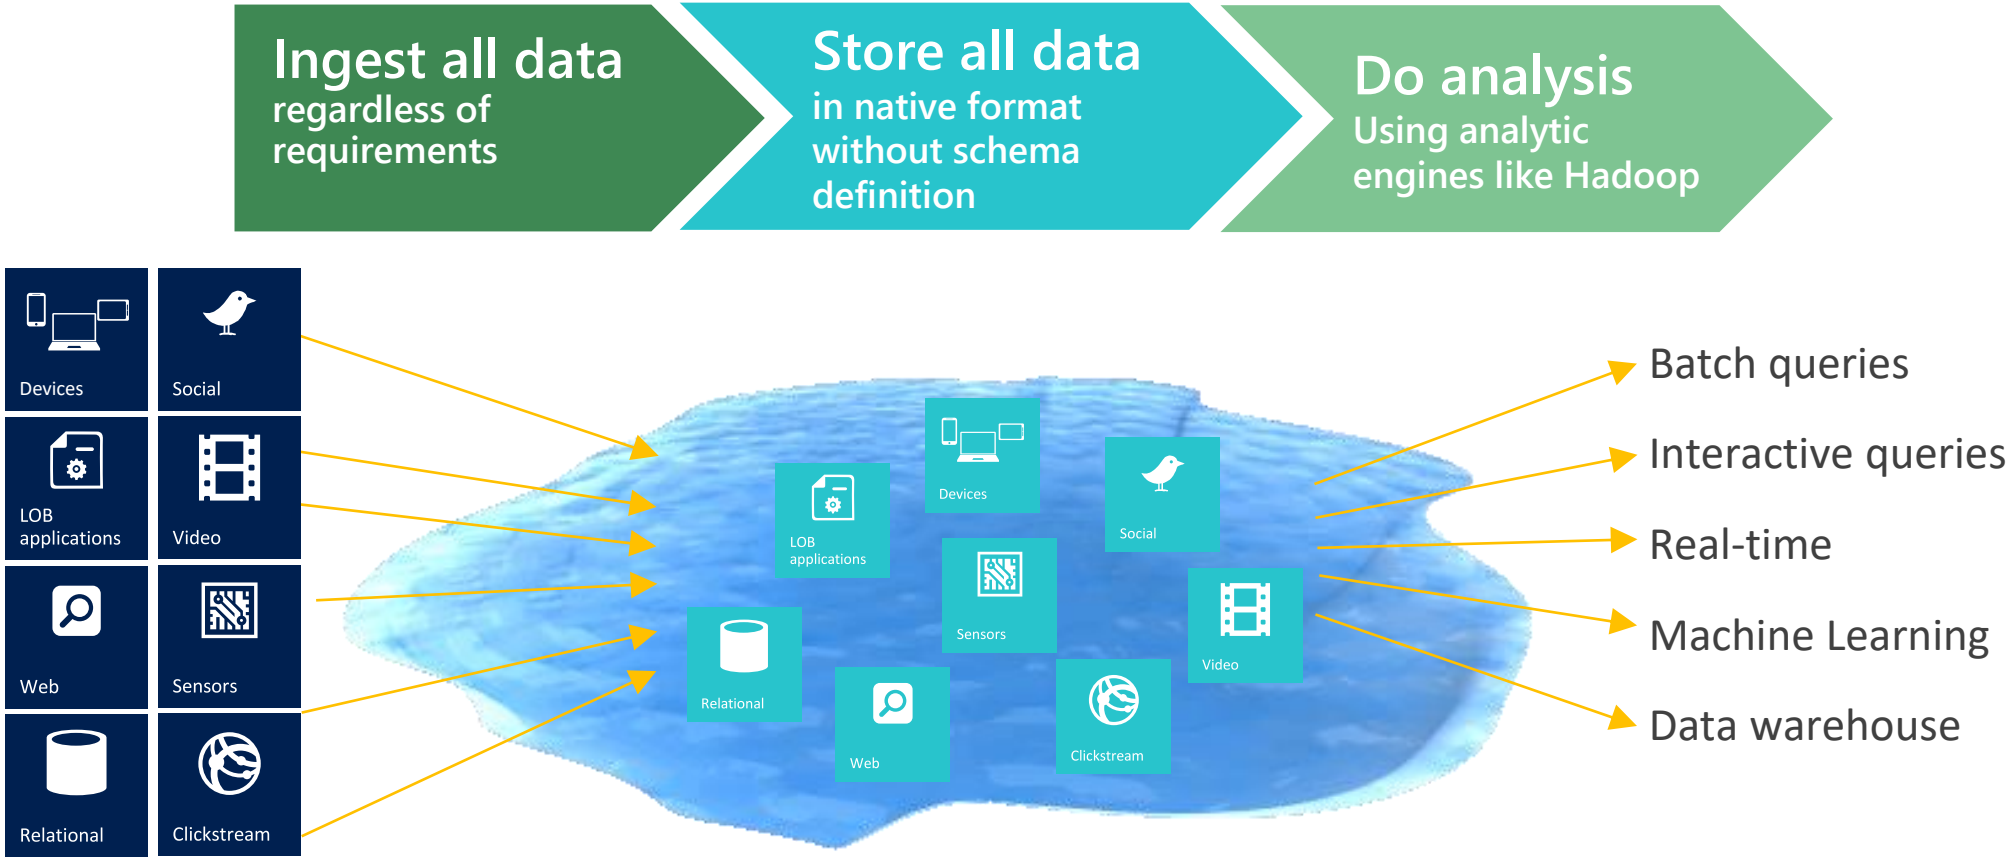

From: M. Rys Presentation

Azure Data Lake

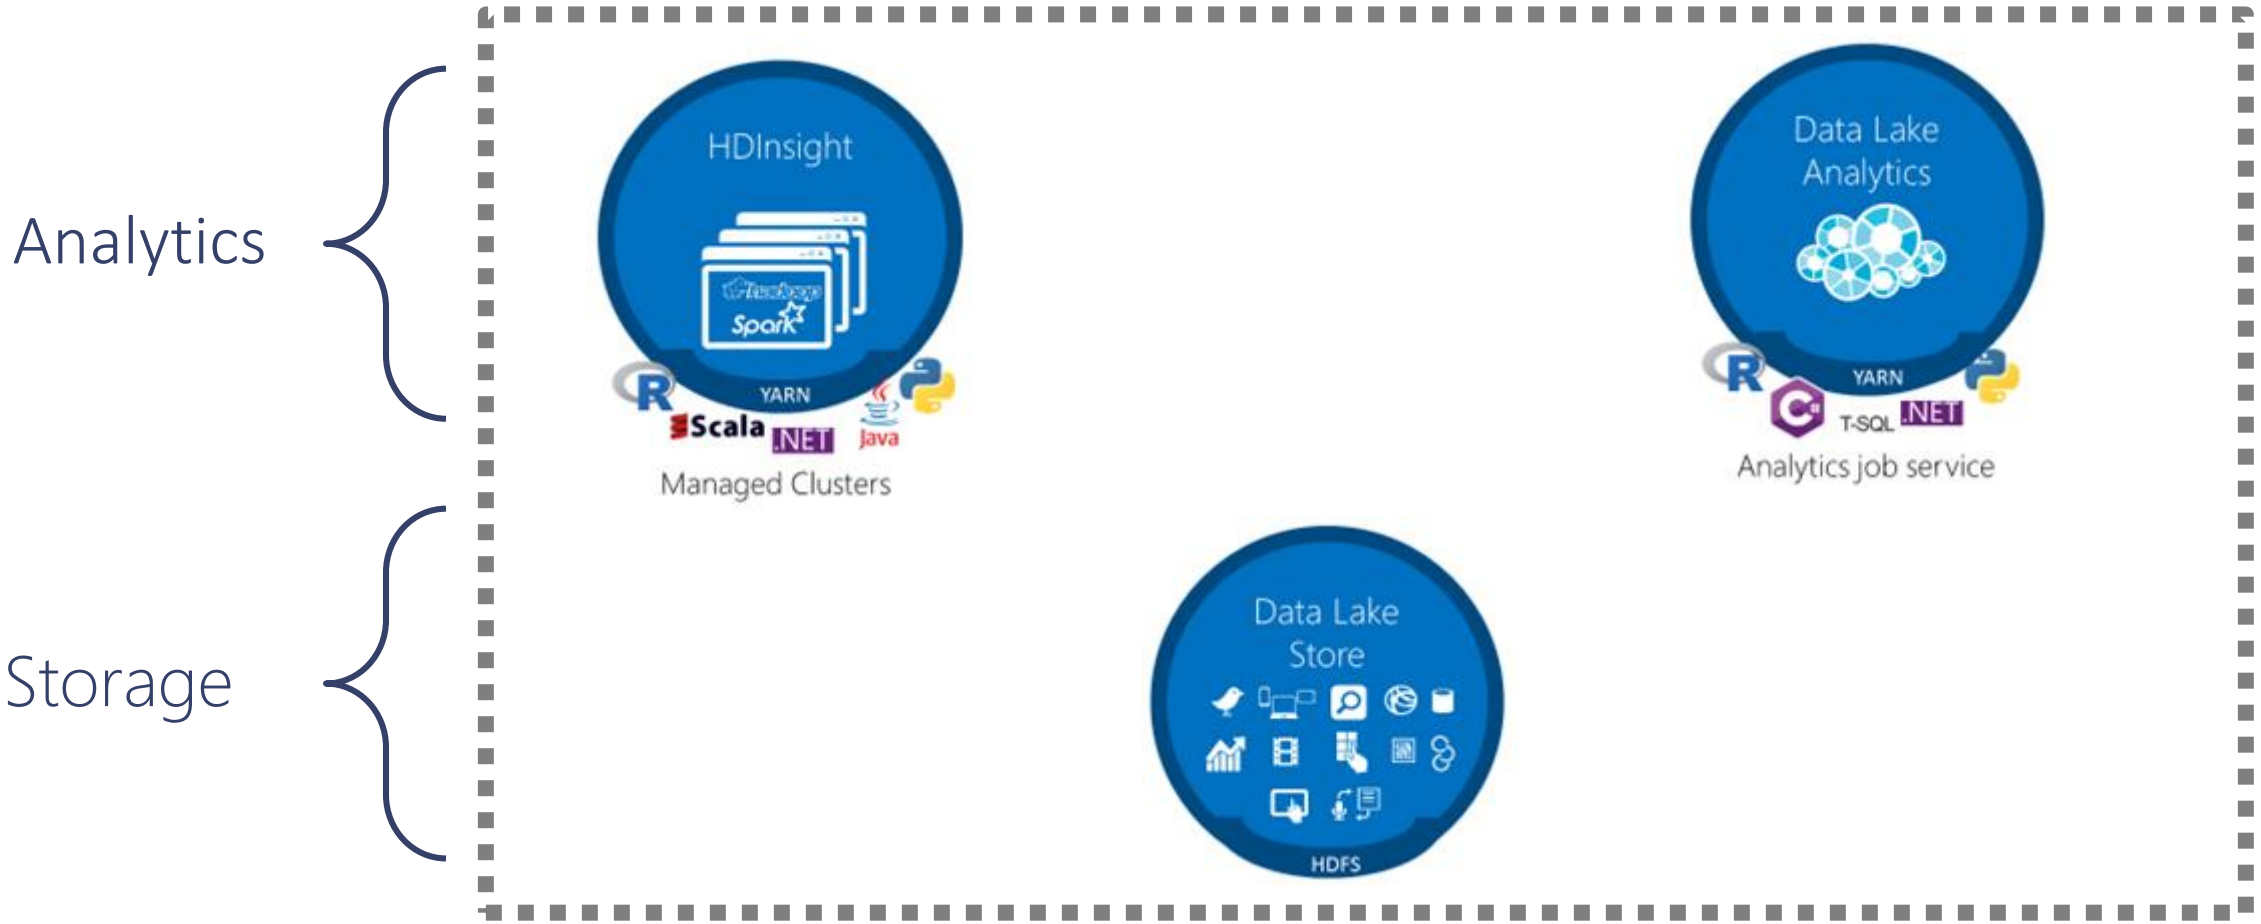

Azure Data Lake

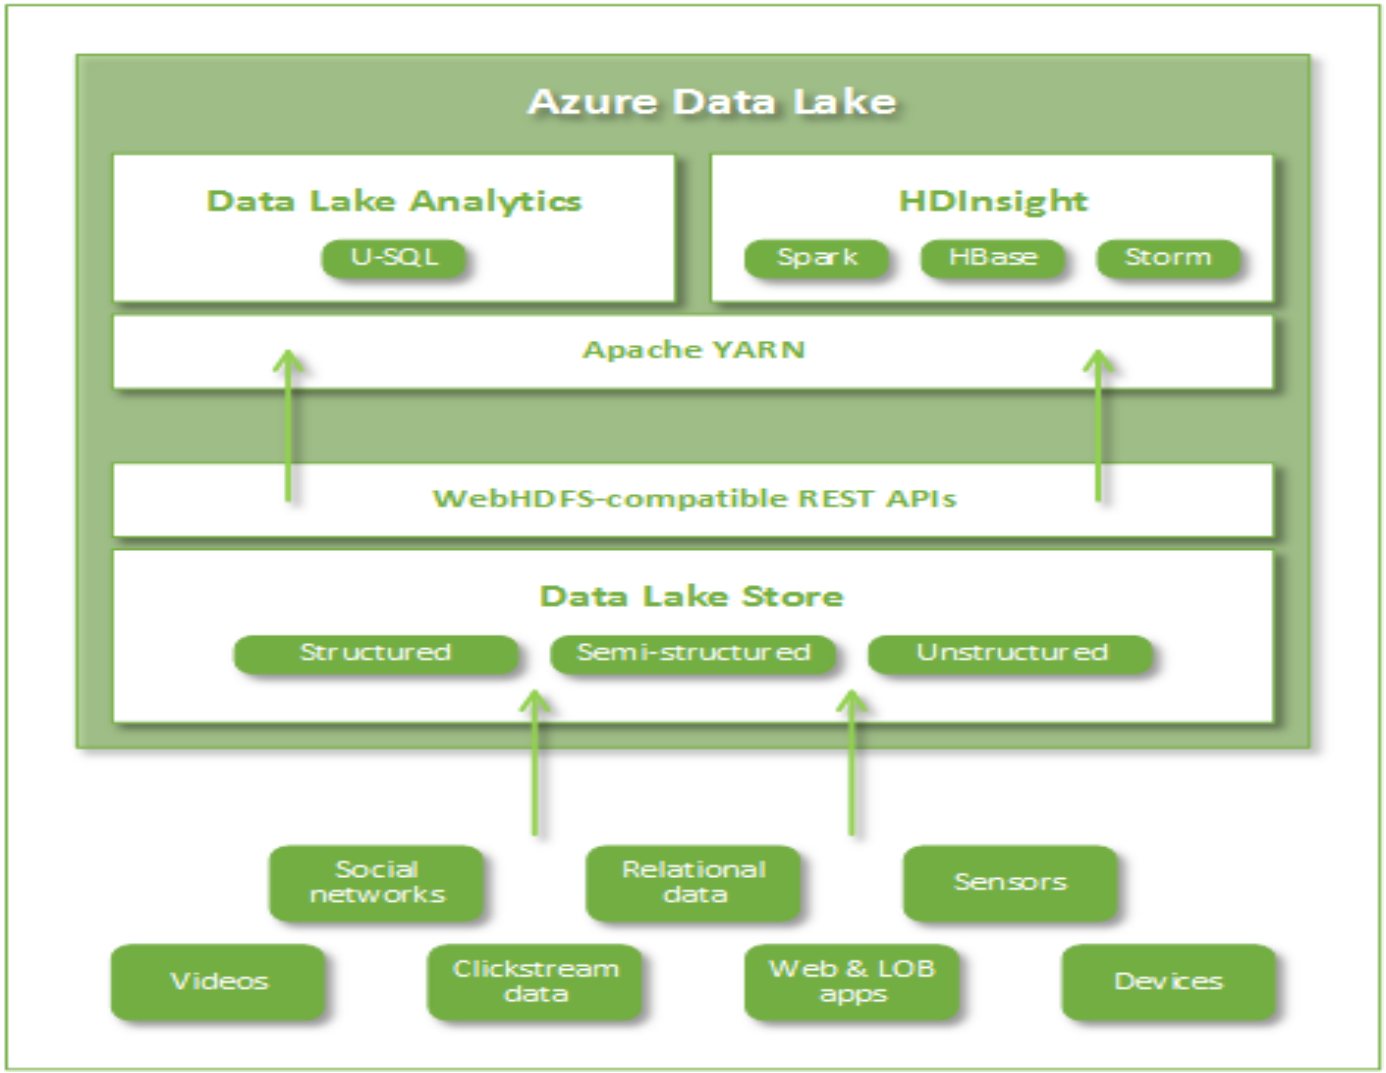

Microsoft

Overview of Azure Data Lake Store

- Built for Hadoop
  - WebHDFS-compatible REST interface
- Unlimited storage, petabyte files
- Performance-tuned for big data analytics
- Highly-available and secure
- Integrates with HDInsight, Cloudera, Hortonworks
- Supports files and folders objects

ADL Store Basics

- Files are split apart into Extents (250 MB)
- For availability and reliability, extents are replicated (3 copies).
- Enables:
  - Parallel read
  - Parallel write

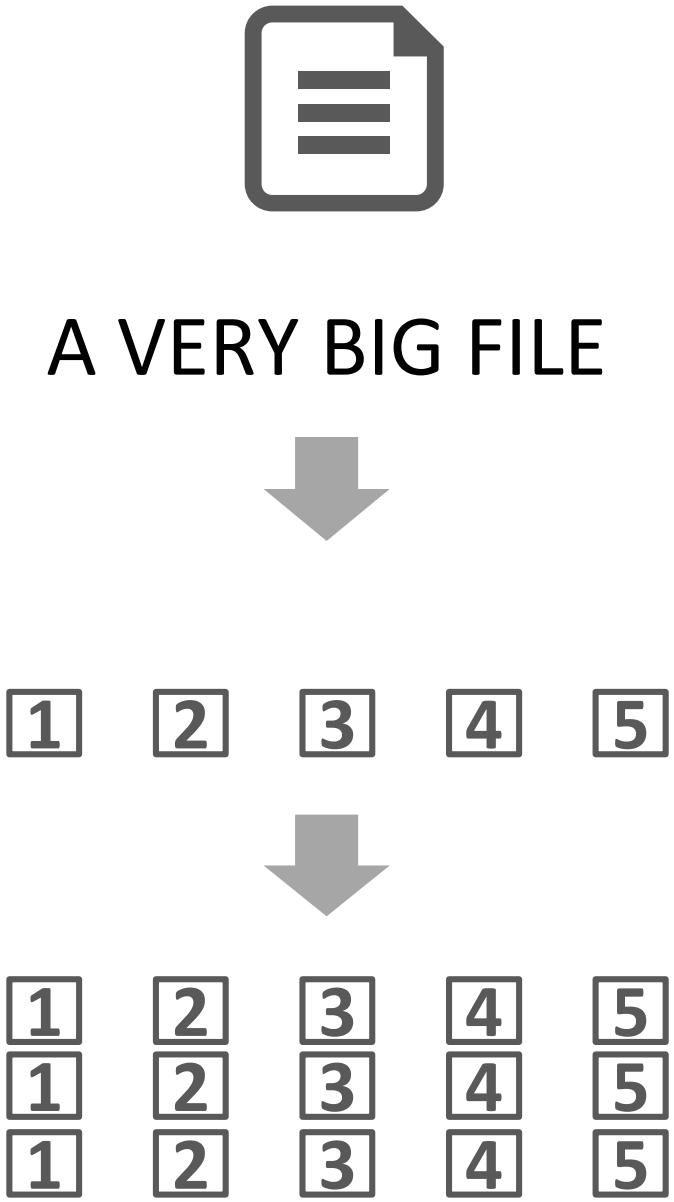

From: S. Reddy Presentation

Working with Azure Data Lake Store

- Local computer
  - Azure Portal
  - Azure PowerShell
  - Azure Cross-platform CLI
  - Using Data Lake Tools for Visual Studio
- Azure Storage Blob
  - Azure Data Factory
  - AdlCopy tool
  - DistCp running on HDInsight cluster
- Streamed data
  - Azure Stream Analytics
  - Azure HDInsight Storm
- Relational data
  - Apache Sqoop
  - Azure Data Factory

Azure Data Explorer

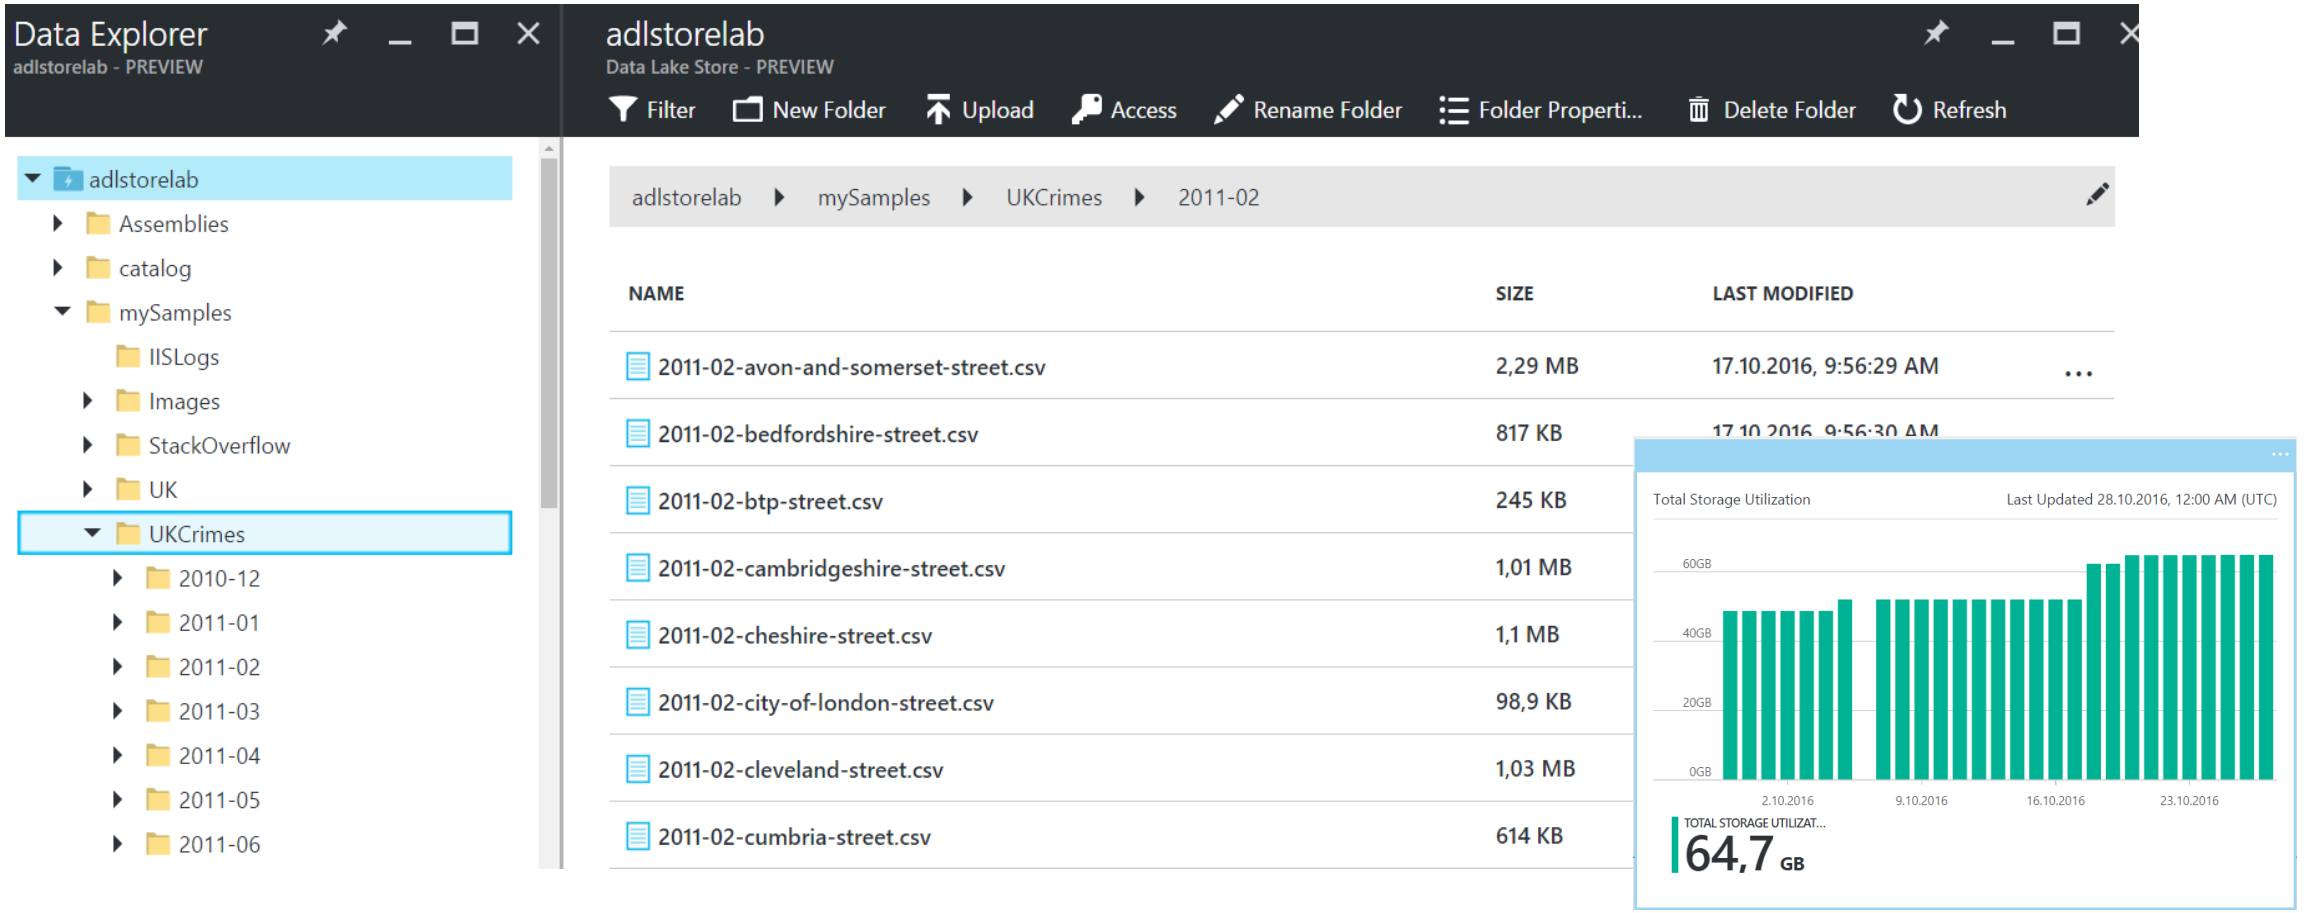

Azure Data Lake Store Pricing

Storage Prices

Storage is available in Pay-as-you-Go and monthly commitment packages.

Pay-as-You-Go

USAGE	PRICE /MONTH
First 100 TB	€0.0329 per GB
Next 100 TB to 1,000 TB	€0.032 per GB
Next 1,000 TB to 5,000 TB	€0.0312 per GB
Over 5,000 TB	Contact Us

Transaction Prices

The following prices apply to transactions performed against your data. The same transaction rates apply for both Pay-as-You-Go as well as Monthly Commitment Packages.

USAGE	PRICE
Write operations (per 10,000)	€0.0422
Read operations (per 10,000)	€0.0034
Delete operations	Free

Azure Data Lake Analytics

A distributed analytics service built on Apache YARN that dynamically scales to your needs

Pay **PER QUERY** & Scale **PER QUERY**

FEDERATED QUERY across Azure data sources

Includes **U-SQL**, a language that unifies the benefits of SQL with the expressive power of C#

No limits to **SCALE**

Optimized to work with **ADL STORE**

From: S. Reddy Presentation

Work across all your cloud Data

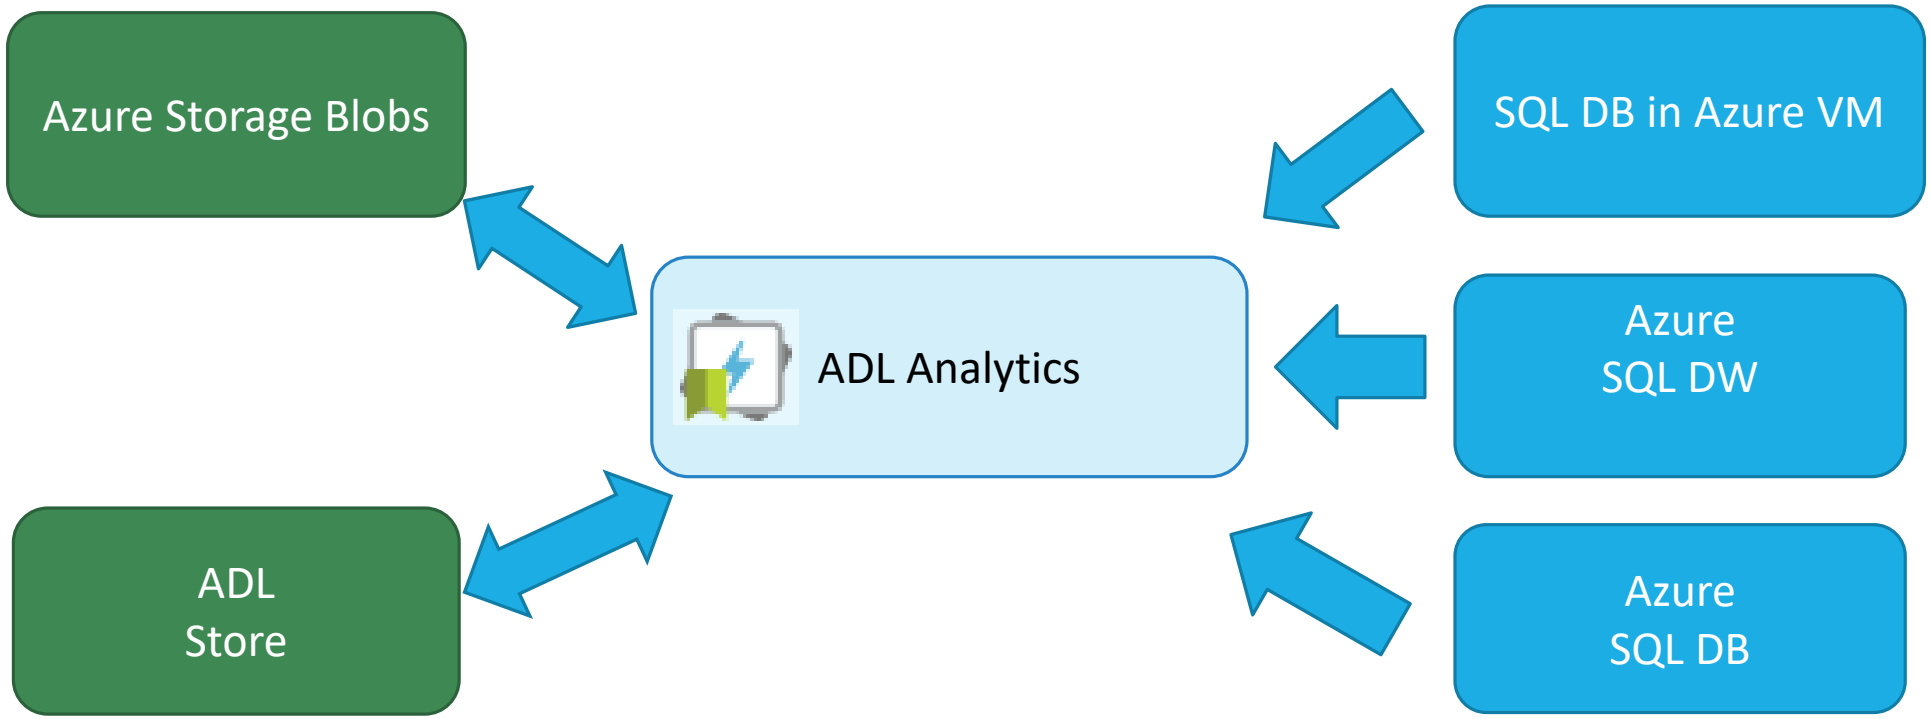

Data Lake Analytics Pricing

Pricing Details

Pay-as-You-Go:

Pay-as-You-Go lets you pay by the second with no long-term commitments.

USAGE	PREVIEW PRICE (UNTIL DECEMBER 31ST, 2016)	GA PRICE (STARTING JANUARY 1ST, 2017)
Analytics Unit	€0.8433/hr	€1.6866/hr
Completed Job	€0.0211 / Job	Free

$$JobCost = (seconds \times ADLU)/3600$$ ***+ Completed Job Cost + Data Lake Transactions Cost***

Azure Data Lake Analytics Unit

Parallelism N = N ADLAUs

1 ADLAU ~=

- A VM with 2 cores and 6 GB of memory

Data Lake Tools for Visual Studio

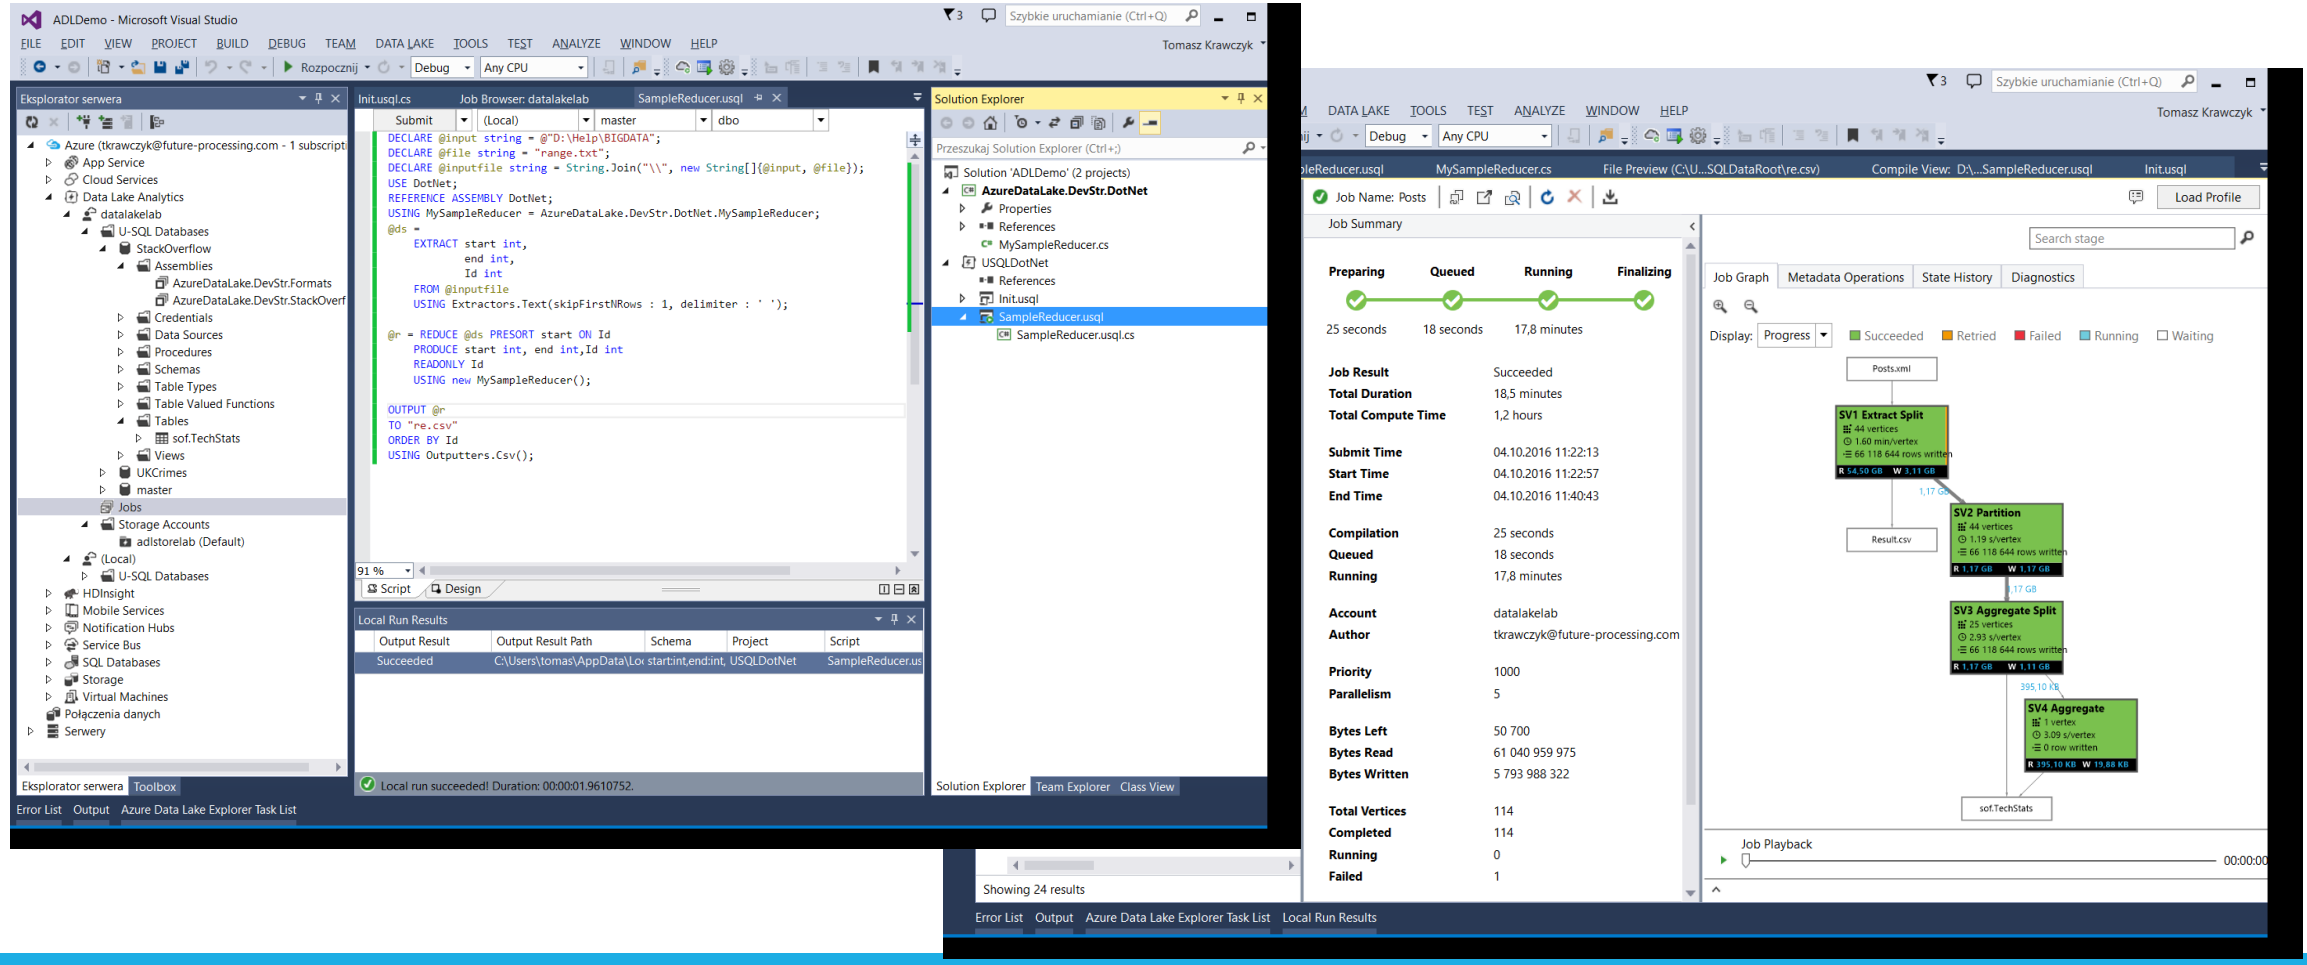

DINNER

U –SQL
A new language for Big Data

Familiar syntax to millions of SQL & .NET developers

Unifies declarative nature of SQL with the imperative power of C#

Unifies structured, semi-structured and unstructured data

Distributed query support over all data

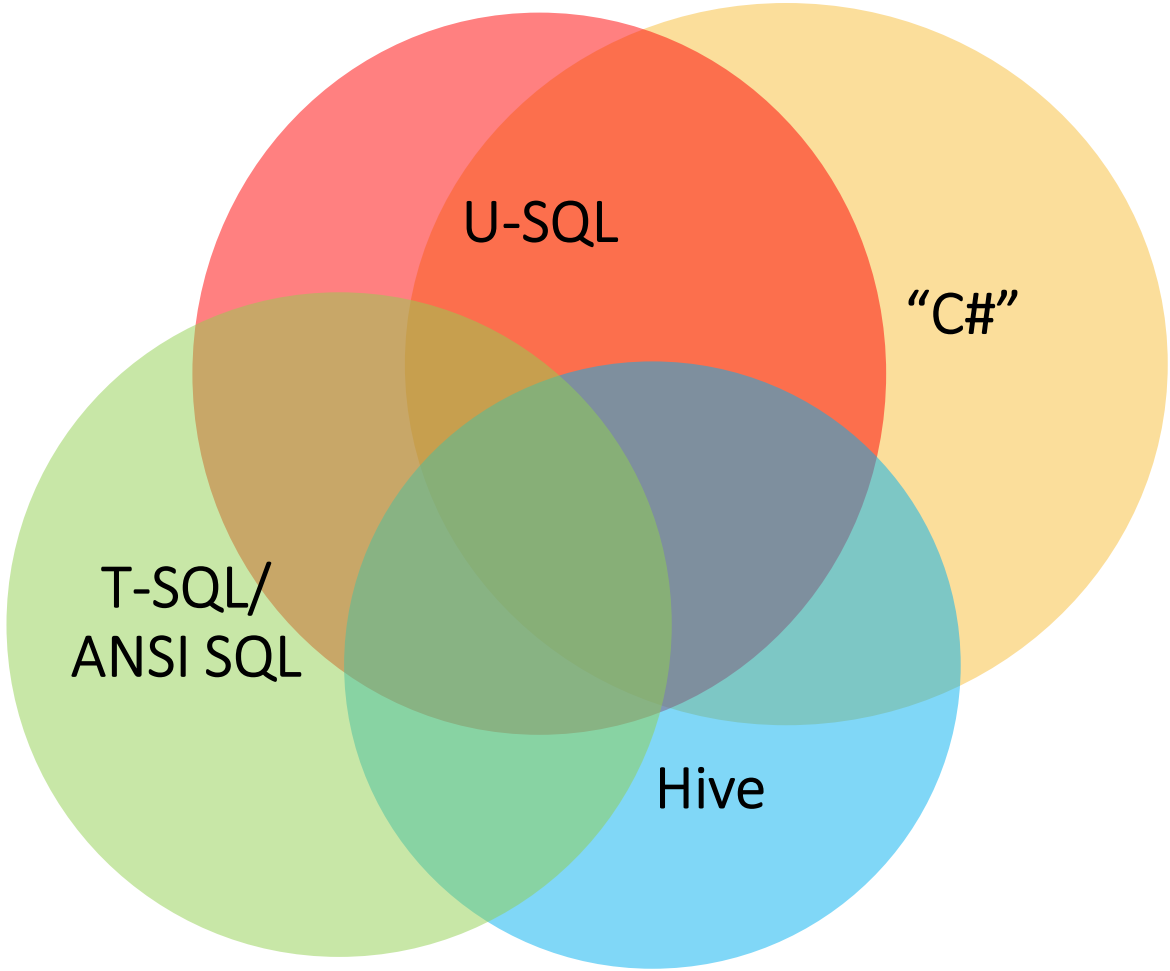

From: M. Rys Presentation

U-SQL SCRIPT

```
DECLARE @inputPostCodes string =
@"mySamples/UK/ukpostcodes.csv";

@postCodes =
    EXTRACT id string,
            postcode string,
            latitude string,
            longitude string
    FROM @inputPostCodes
    USING Extractors.Csv(skipFirstNRows : 1);

OUTPUT @postCodes
TO @"ukpostcodes.txt"
USING Outputters.Text();
```

DECLARE (Optional)

EXTRACT (or SELECT)

Apply Schema on read

OUTPUT (or INSERT)

U-SQL Meta Data Catalog

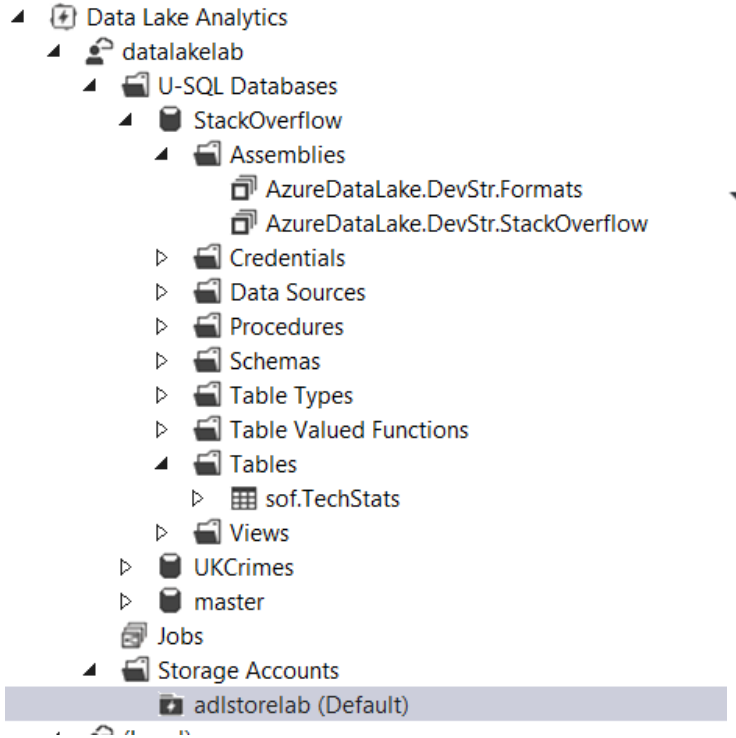

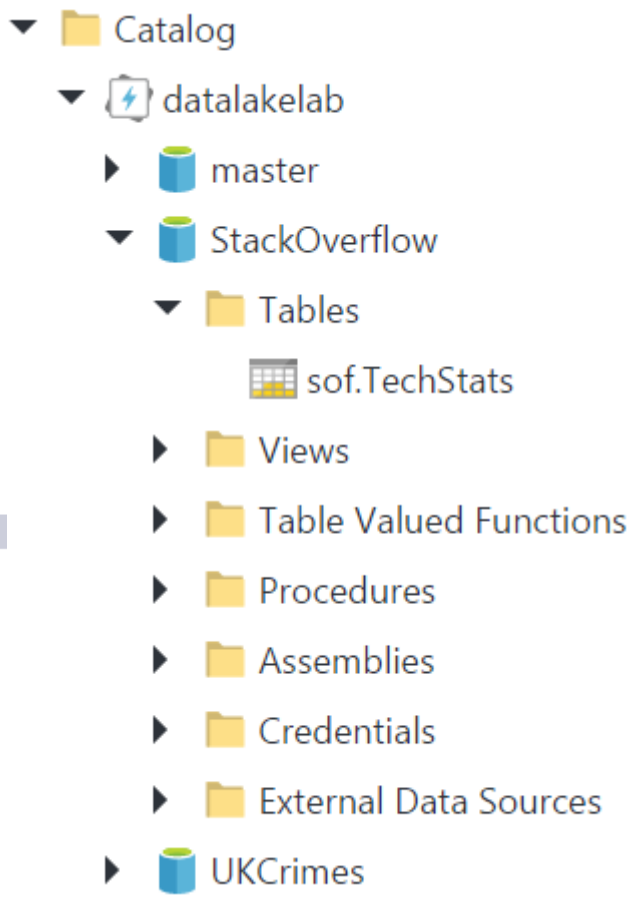

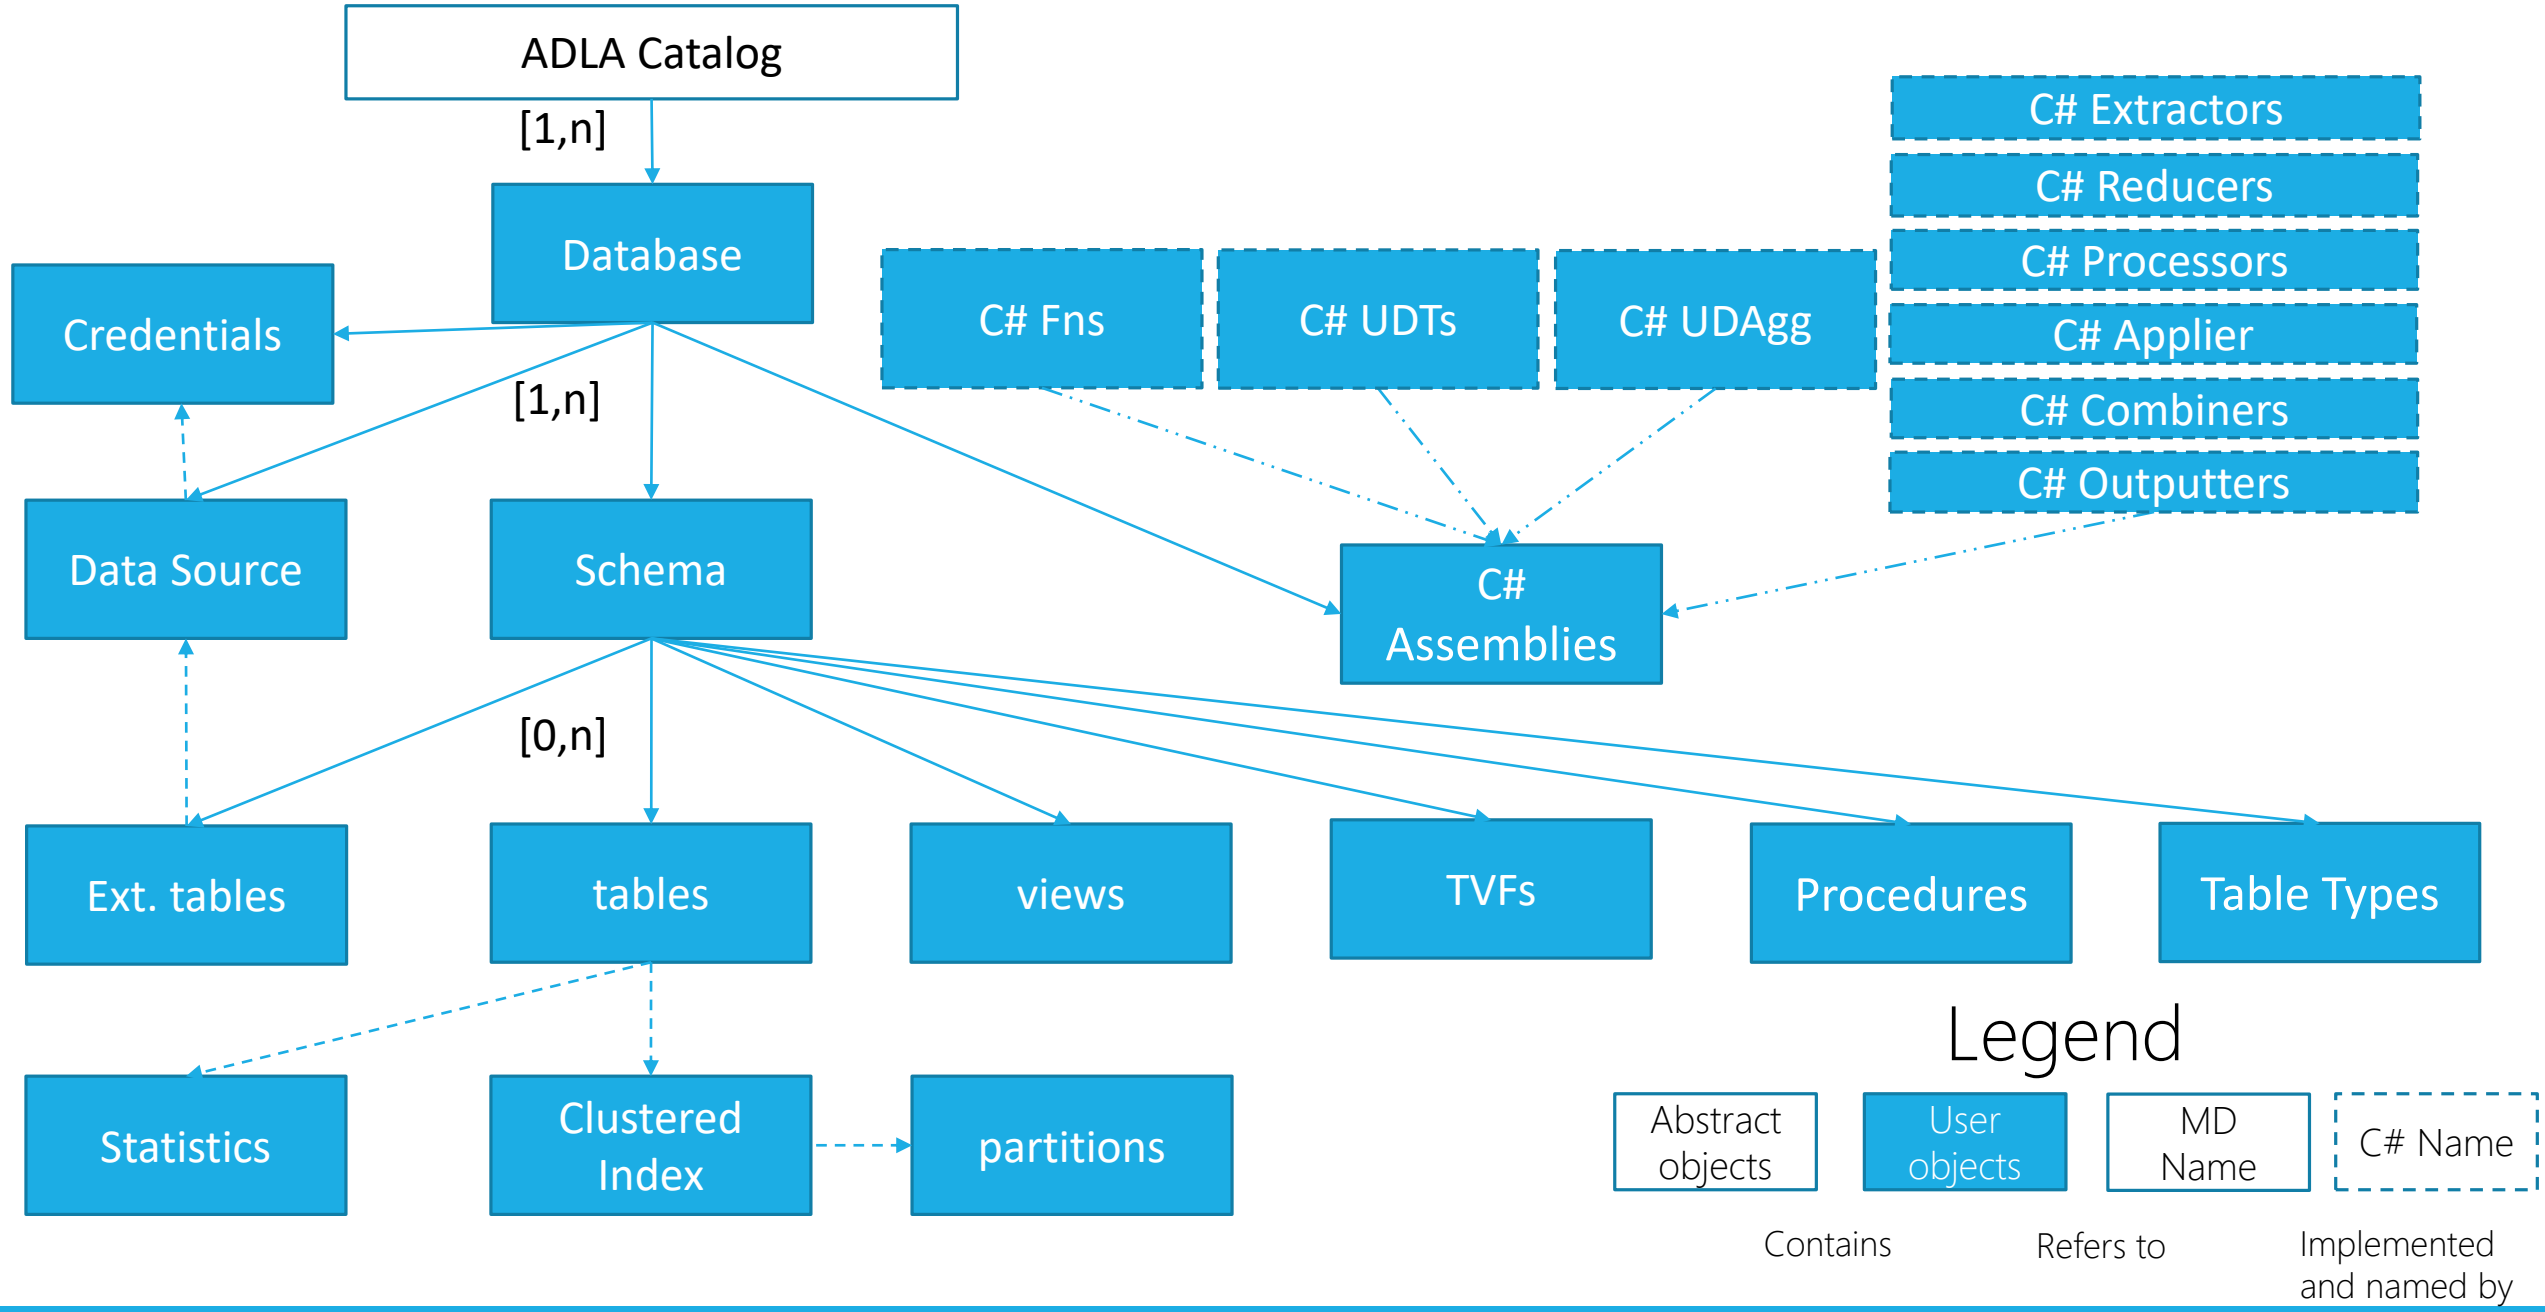
ADLA Catalog
[1,n]
Database
Credentials
[1,n]
Data Source
Schema
C# Fns
C# UDTs
C# UDAgg
C# Extractors
C# Reducers
C# Processors
C# Applier
C# Combiners
C# Outputters
C# Assemblies
[0,n]
Ext. tables
tables
views
TVFs
Procedures
Table Types
Statistics
Clustered Index
partitions
Legend
Abstract objects
User objects
MD Name
C# Name
Contains
Refers to
Implemented and named by

U-SQL DECLARE VARIABLES

```
DECLARE @text1 string = "Columbus Day";
DECLARE @text2 string = @"Columbus Day";
DECLARE @text3 char = 'a';
DECLARE @text4 string = "BEGIN" + @text1 + "END";
DECLARE @text5 string = string.Format("BEGIN{0}END", @text1);
DECLARE @text6 string = string.Join(" ", new String[]{@text1, "4.0"});

DECLARE @numeric1 sbyte = 0;
DECLARE @numeric2 short = 1;
DECLARE @numeric3 int = 2;
DECLARE @numeric4 long = 3L;
DECLARE @numeric5 float = 4.0f;
DECLARE @numeric6 double = 5.0;

DECLARE @d1 DateTime = System.DateTime.Parse("1979/03/31");
DECLARE @d2 DateTime = DateTime.Now;

DECLARE @misc1 bool = true;
DECLARE @misc2 Guid = System.Guid.Parse("BEF7A4E8-F583-4804-9711-7E608215EBA6");
DECLARE @misc4 byte [] = new byte[] { 0, 1, 2, 3, 4};
```

U-SQL EXTRAXTORS and OUTPUTTERS

- Csv
- Text
- Tsv

```
public static Microsoft.Analytics.Interfaces.IExtractor Csv(System.Text.Encoding encoding);
public static Microsoft.Analytics.Interfaces.IExtractor Csv(System.String rowDelimiter,
    System.Nullable<System.Char> escapeCharacter, System.String nullEscape,
    System.Text.Encoding encoding, System.Boolean quoting, System.Boolean silent,
    System.Int32 skipFirstNRows, System.String charFormat);
public static Microsoft.Analytics.Interfaces.IExtractor Text(System.Text.Encoding encoding);
public static Microsoft.Analytics.Interfaces.IExtractor Text(System.Char delimiter,
    System.String rowDelimiter, System.Nullable<System.Char> escapeCharacter, System.String nullEscape,
    System.Text.Encoding encoding, System.Boolean quoting, System.Boolean silent, System.Int32 skipFirstNRows,
    System.String charFormat);
public static Microsoft.Analytics.Interfaces.IExtractor Tsv(System.Text.Encoding encoding);
public static Microsoft.Analytics.Interfaces.IExtractor Tsv(System.String rowDelimiter,
    System.Nullable<System.Char> escapeCharacter, System.String nullEscape,
    System.Text.Encoding encoding, System.Boolean quoting, System.Boolean silent,
    System.Int32 skipFirstNRows, System.String charFormat);
```

U-SQL FILESETS

```
DECLARE @inputCrimes = @"mySamples/UKCrimes/{Date:yyyy}-
{Date:MM}/{Input}-street.csv";
@crimes =
    EXTRACT CrimeID string,
            Month string,
            ReportedBy string,
            FallsWithin string,
            Longitude string,
            Latitude string,
            Location string,
            LSOACode string,
            LSOAName string,
            CrimeType string,
            LastOutcomeCategory string,
            Context string,
            Date DateTime,
            Input string
    FROM @inputCrimes
    USING Extractors.Csv(silent :
false,skipFirstNRows:1);
```

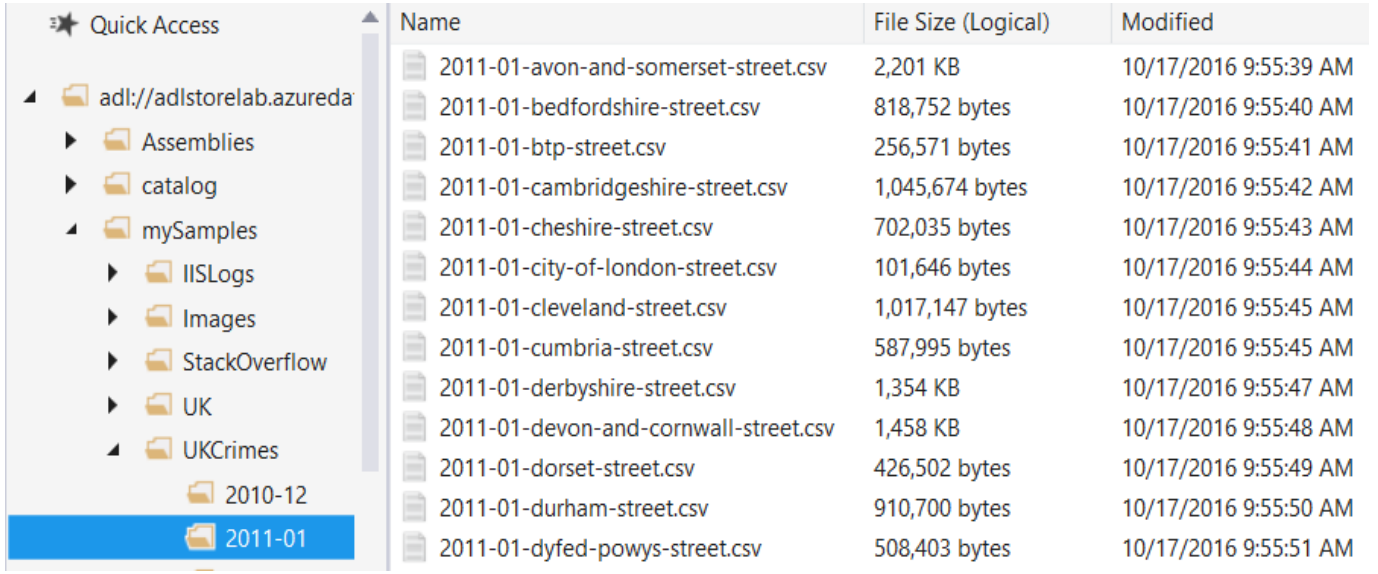

Virtual columns

WORKSHOP 1

Azure Data Lake

https://github.com/rkostrzewski/usql-workshop

U-SQL FILTERING

- ROWSET(s)
- TABLE(s)
- WHERE
- AND & OR
- ==,>=,!= (C# OPERATOR(s))
- CONTAINS (C# string)

```
@crimeInDay =
    SELECT CrimeType,
           Date,
           COUNT( * ) AS Count
    FROM @crimes
    WHERE Date >= @crimeDate AND
(CrimeType.Contains("Other")
    OR
CrimeType.StartsWith("Robb"))
```

U-SQL ROWSETS

```
@postCodes =
    EXTRACT id string,
            postcode string,
            latitude string,
            longitude string
    FROM @inputPostCodes
    USING Extractors.Csv(skipFirstNRows:1);

@topCities =
    EXTRACT id int,
            name string,
            population string,
            postcode string
    FROM @input10topCities
    USING Extractors.Text(delimiter : ';');

@topCitiesWithGPS =
    SELECT tc.name,tc.population, pc.latitude,pc.longitude
    FROM @topCities AS tc
         JOIN
             @postCodes AS pc
         ON pc.postcode == tc.postcode;
```

Rowset

Rowset

Rowset

WORKSHOP 2

U-SQL ARRAY and MAP

SQL.ARRAY<T> == IList<T>

```
@m = SELECT new SqlArray<string>
(
        tweet.Split(
        new char[]{''}).Where(x => x.StartsWith("@")))
        AS mentions
    FROM @t;

@m = SELECT m.Substring(1) AS m
          , "mention" AS category
    FROM @m CROSS APPLY EXPLODE(mentions) AS t(m)
```

SQL.MAP<T,U> ==IDictionary<T,U>

```
@ds =
    SELECT content,fileName, new SQL.MAP<int,string>() AS
colors
    FROM @rs;

@ds =
    PROCESS @ds
    PRODUCE content,colors,fileName
            READONLY fileName
    USING new ImageColorsProcessor(4);

@ds =
    SELECT fileName,
           order,
           colorName
    FROM @ds
         CROSS APPLY
             EXPLODE(colors) AS colors(order, colorName);
```

U-SQL SORTING

ROWSET

```
@distances =
    SELECT CrimeId,
           CityName,
           CrimeType,
           Year,
           Month
    FROM @merged
ORDER BY CityName DESC
FETCH FIRST 10 ROWS;
```

OUPUT

```
OUTPUT @ds
TO "result.csv"
ORDER BY fileName,
         order
USING Outputters.Csv();
```

SELECT with ORDER BY requires a FETCH FIRST

WORKSHOP 3

USQL -AGGREGATIONS

GROUP BY

HAVING

AGGREGATIONS

- MAX
- MIN
- SUM
- MAX
- MIN
- SUM
- **ARRAY_AGG**

```
@output =
    SELECT
        MAX(Duration) AS DurationMax,
        MIN(Duration) AS DurationMin,
        AVG(Duration) AS DurationAvg,
        SUM(Duration) AS DurationSum,
        VAR(Duration) AS DurationVarianve,
        STDEV(Duration) AS DurationStDev,
    FROM @searchlog
    GROUP BY Region
    HAVING DurationMin > 1;
```

U-SQL WINDOW FUNCTIONS

RANKING FUNCTIONS

- RANK
- DENSE_RANK
- NTILE
- ROW_NUMBER

ANALIYTIC WINDOW FUNCTIONS

- CUME_DIST
- PERCENT_RANK
- PERCENTILE_CONT
- PERCENTILE_DISC
- CUME_DIST

```
@result =
SELECT
    *,
    ROW_NUMBER() OVER (PARTITION BY Vertical ORDER BY
Latency) AS RowNumber,
    RANK() OVER (PARTITION BY Vertical ORDER BY Latency)
AS Rank,
    DENSE_RANK() OVER (PARTITION BY Vertical ORDER BY
Latency) AS DenseRank
FROM @querylog;
```

WORKSHOP 4

U-SQL DATABASES AND SCHEMES

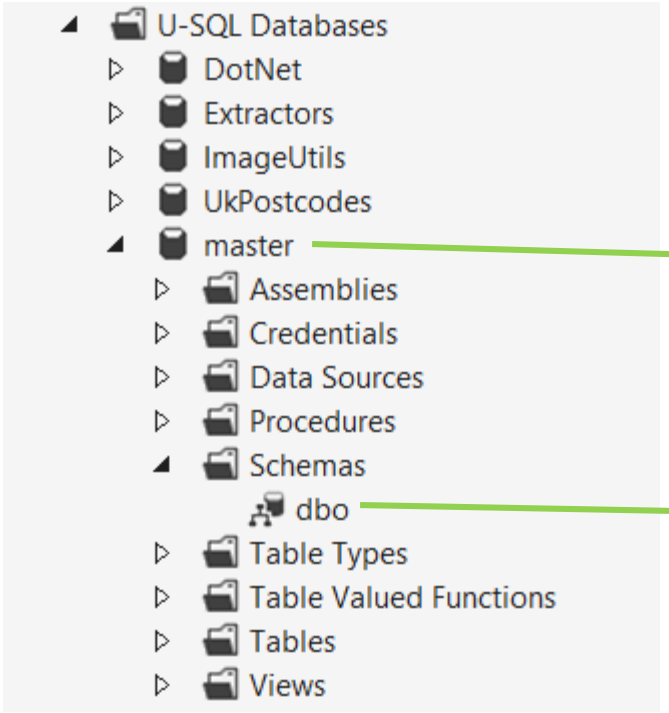


```
CREATE DATABASE IF NOT EXISTS UKCrimes;
USE DATABASE UKCrimes;
CREATE SCHEMA IF NOT EXISTS cr;
```

Default database

Default schema

U-SQL TABLES

- MANAGED TABLES and EXTERNAL TABLES
- ONLY INSERT
- CONSISTS OF FOUR THINGS:
  - A NAME
  - COLUMNS
  - A CLUSTERED INDEX
  - PARTITIONING SCHEME

```
DROP TABLE IF EXISTS vehiclesP;
CREATE TABLE vehiclesP(
      vehicle_id int
    , entry_id long
    , event_date DateTime
    , latitude float
    , longitude float
    , speed int
    , direction string
    , trip_id int?
    , INDEX idx CLUSTERED (vehicle_id ASC)
      PARTITIONED BY (event_date)
      DISTRIBUTED BY HASH (vehicle_id) INTO 4
);
```

U-SQL VIEWS and FUNCTIONS

VIEWS

```
CREATE VIEW IF NOT EXISTS vCrimes
    AS
EXTRACT CrimeID string,
        Month string,
        Date DateTime,
        Input string
FROM @"\UKCrimesCities\{Date:yyyy}-{Date:MM}\{Input}-street.csv"
USING Extractors.Csv(silent : false,
 skipFirstNRows : 1);
```

FUNCTIONS (TVF)

```
CREATE FUNCTION tvf_Crimes(@input string)
RETURNS @result TABLE(CrimeID string,
 Month string)
AS
BEGIN
    @crimes =
    EXTRACT CrimeID string,
            Month string
    FROM @input
    USING Extractors.Csv(silent : false,
skipFirstNRows:1);

    @result = SELECT CrimeID,
            Month
            Input FROM @crimes;
    END;
```

U-SQL JOINS

- INNER JOIN
- FULL OUTER JOIN
- LEFT OUTER JOIN
- RIGHT OUTER JOIN
- CROSS JOIN
- LEFT SEMIJOIN (IN)
- RIGHT SEMIJOIN (IN)
- LEFT ANTISEMIJOIN (NOT IN)
- RIGHT ANTISEMIJOIN (NOT IN)

```
@topCitiesWithGPS =

    SELECT tc.name,tc.population,
pc.latitude,pc.longitude

    FROM @topCities AS tc

         JOIN

             @postCodes AS pc

         ON pc.postcode ==
tc.postcode;
```

U-SQL C# METHODS

```
@distances =

    SELECT CrimeId,

           CityName,

           CrimeType,

           Year,

           Month,

           Gps.ComputeDistance

           (sLatitude, sLongitude, dLatitude, dLongitude)
AS Distance

    FROM @merged;
```

```
public static double ComputeDistance(double sLat, double
sLong, double dLat, double dLong)

        {

            var locA = new GeoCoordinate(sLat, sLong);

            var locB = new GeoCoordinate(dLat, dLong);

            return locA.GetDistanceTo(locB); // metres

        }
```

C# Method

IExtractor

```
[SqlUserDefinedExtractor(AtomicFileProcessing = true)]
    public class BinaryContentExtractor : IExtractor
    {
        public override IEnumerable<IRow> Extract(IUnstructuredReader input, IUpdatableRow output)
        {
            using (var ms = new MemoryStream())
            {
                input.BaseStream.CopyTo(ms);
                var content = ms.ToArray();
                output.Set(0, content);
                yield return output.AsReadOnly();
            }
        }
    }
```

U-SQL USING ASSEMBLIES

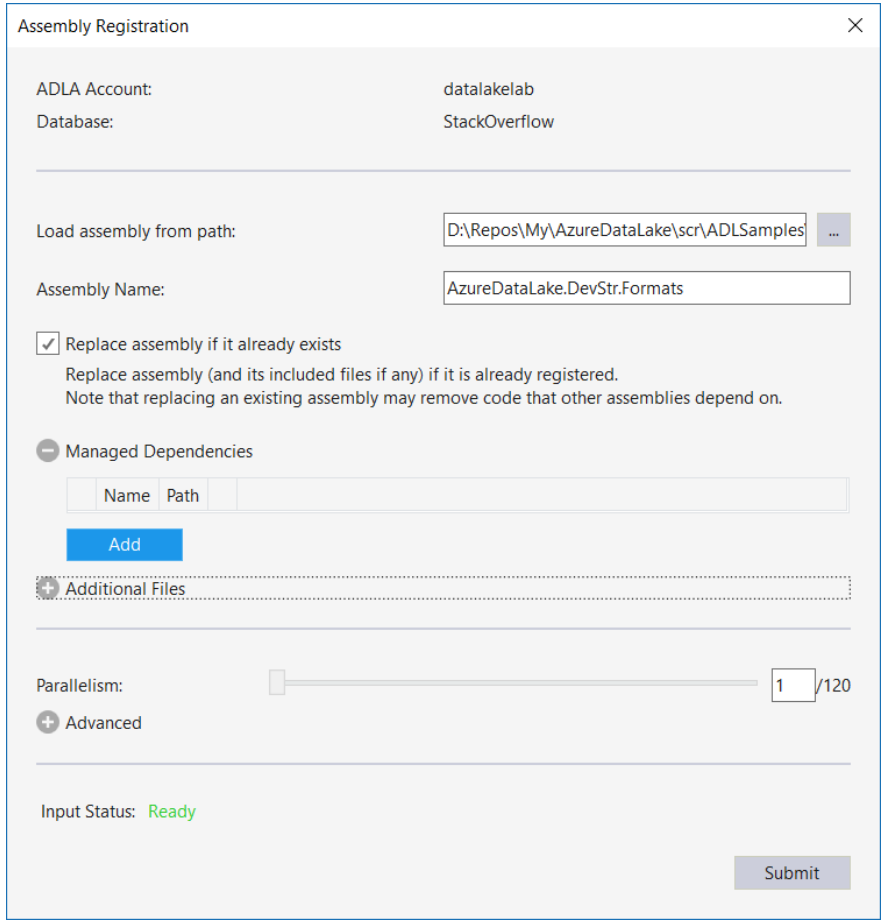


```
DECLARE @myAssemblyPath string =
@"D:\Repos\AzureDataLake.DevStr.Formats\bin\Debug\";

DECLARE @myAssemblyName string =
@myAssemblyPath+"AzureDataLake.DevStr.Formats.dll";

CREATE DATABASE IF NOT EXISTS Extractors;

USE DATABASE Extractors;

DROP ASSEMBLY IF EXISTS MyExtractors;

CREATE ASSEMBLY MyExtractors

FROM @myAssemblyName;
```

U-SQL USING ASSEMBLIES

```
DECLARE @imgFile string = @"D:\Help\BIGDATA\Images\{fileName}.jpg";
USE DATABASE Extractors;
REFERENCE ASSEMBLY MyExtractors;
USING BinaryExtractor = AzureDataLake.DevStr.Formats.BinaryContentExtractor;
REFERENCE ASSEMBLY ImageUtils.ImageUtils;
USING ImageColorsProcessor = AzureDataLake.DevStr.ImageUtils.ImageColorProducer;
@rs =
    EXTRACT content byte[],
            fileName string
    FROM @imgFile
    USING new BinaryExtractor();

@ds =
    SELECT content,fileName, new SQL.MAP<int,string>() AS colors
    FROM @rs;

@ds =
    PROCESS @ds
    PRODUCE content,colors,fileName
            READONLY fileName
    USING new ImageColorsProcessor(4);
```

Reference

Alias

External Extractor

External Processor

U-SQL

WORKSHOP 5

COFFEE BREAK

Azure Data Lake Analytics Jobs Execution

U-SQL Script -> Job Graph Logical -> Physical Plan

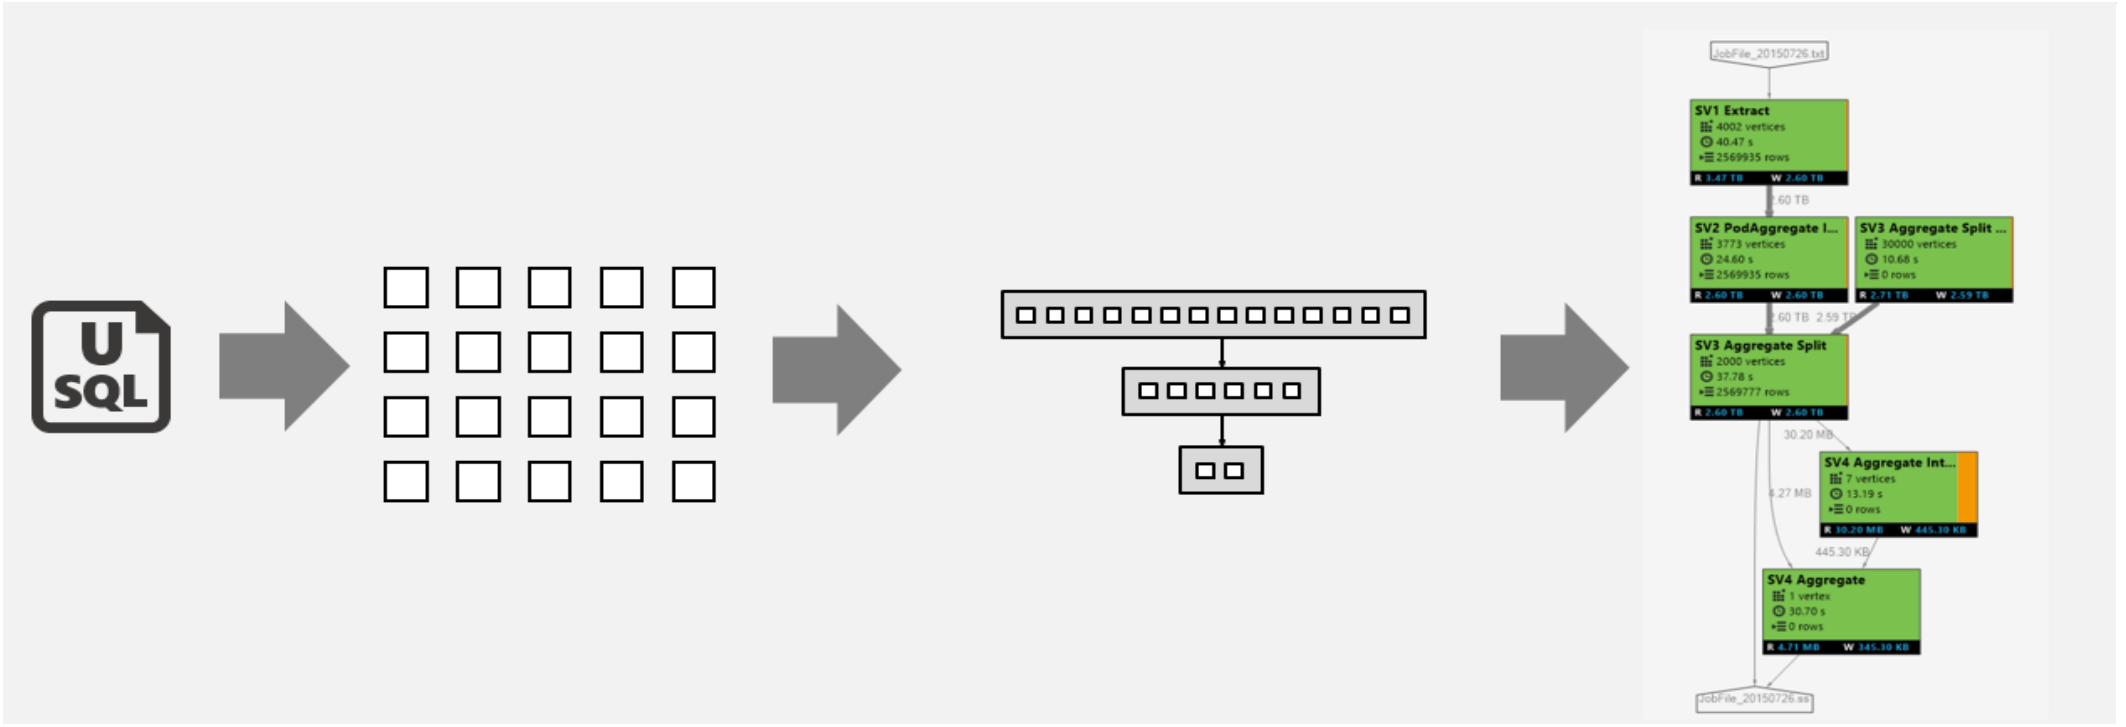

Each square = "a vertex" represents a fraction of the total

Vertexes in each SuperVertex (aka "Stage) are doing the same operation on a different part of the same data.

Visualized as a "Job Graph"

From: S. Reddy Presentation

Azure Data Lake Analytics Jobs

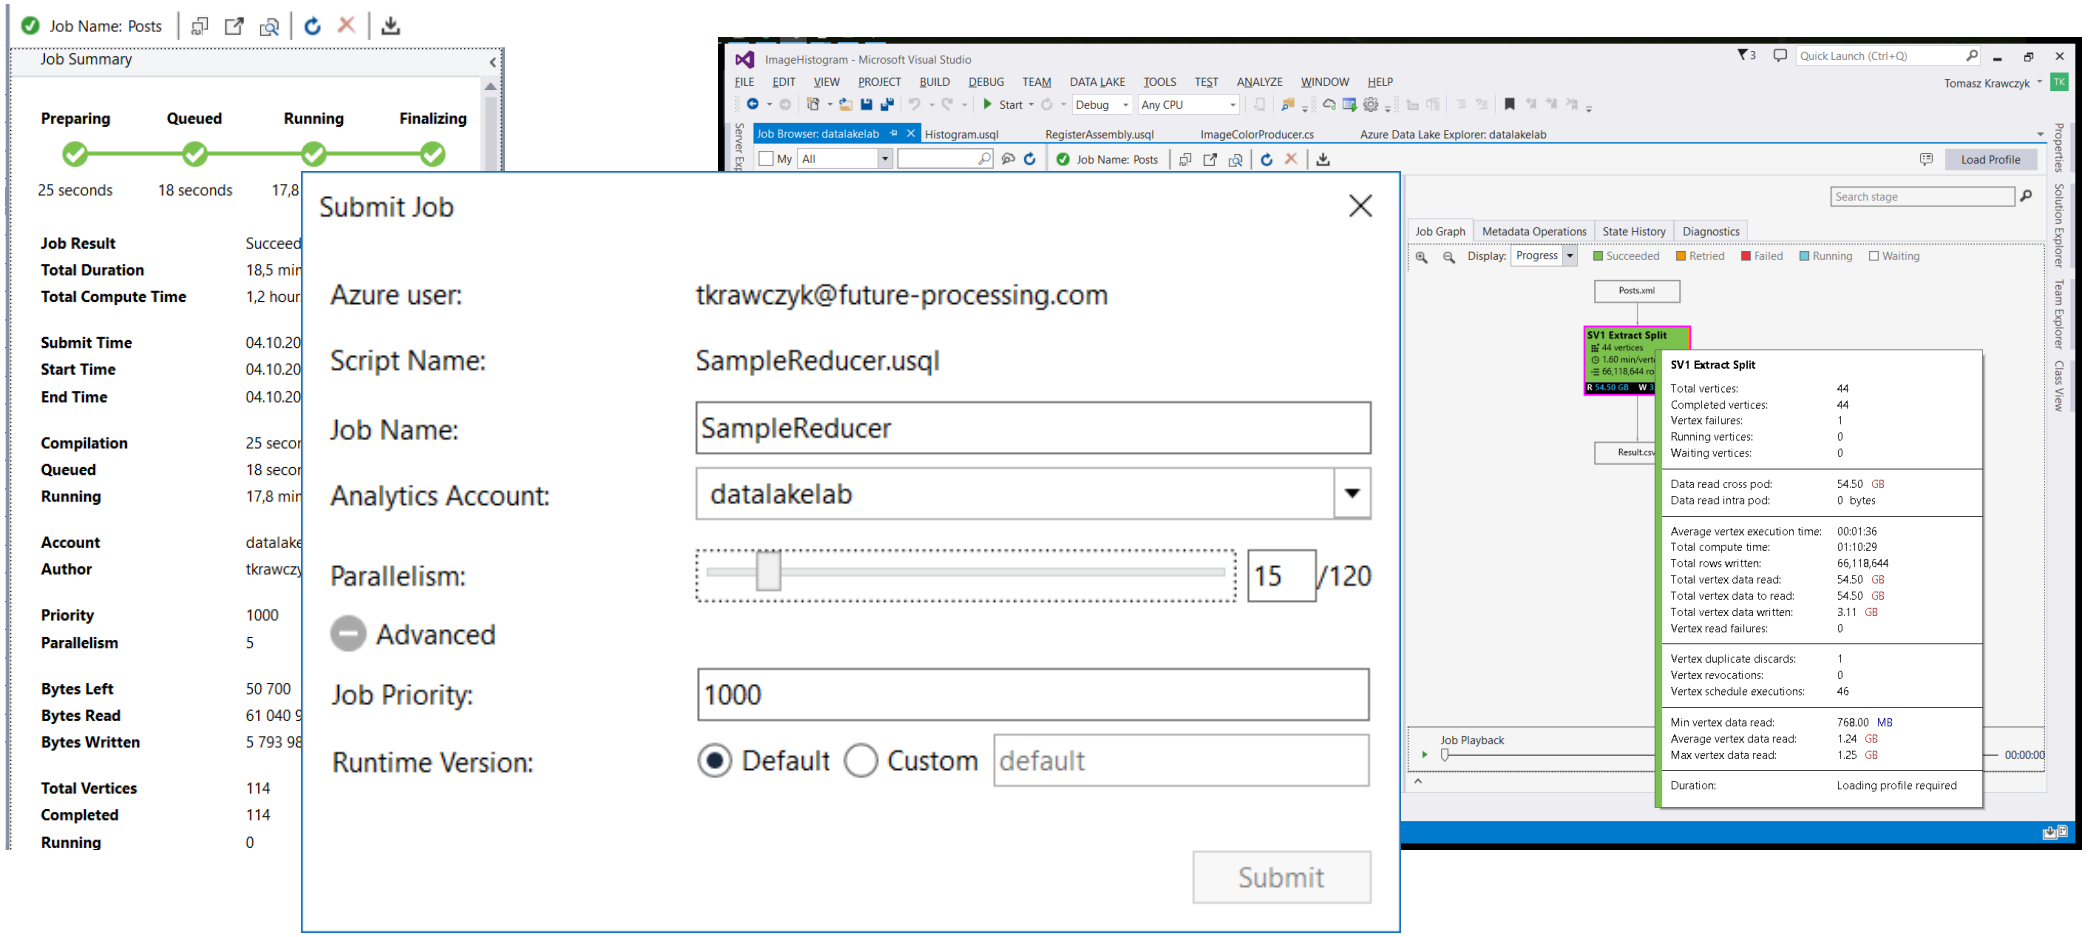

Data Lake Analytics - AU Usage

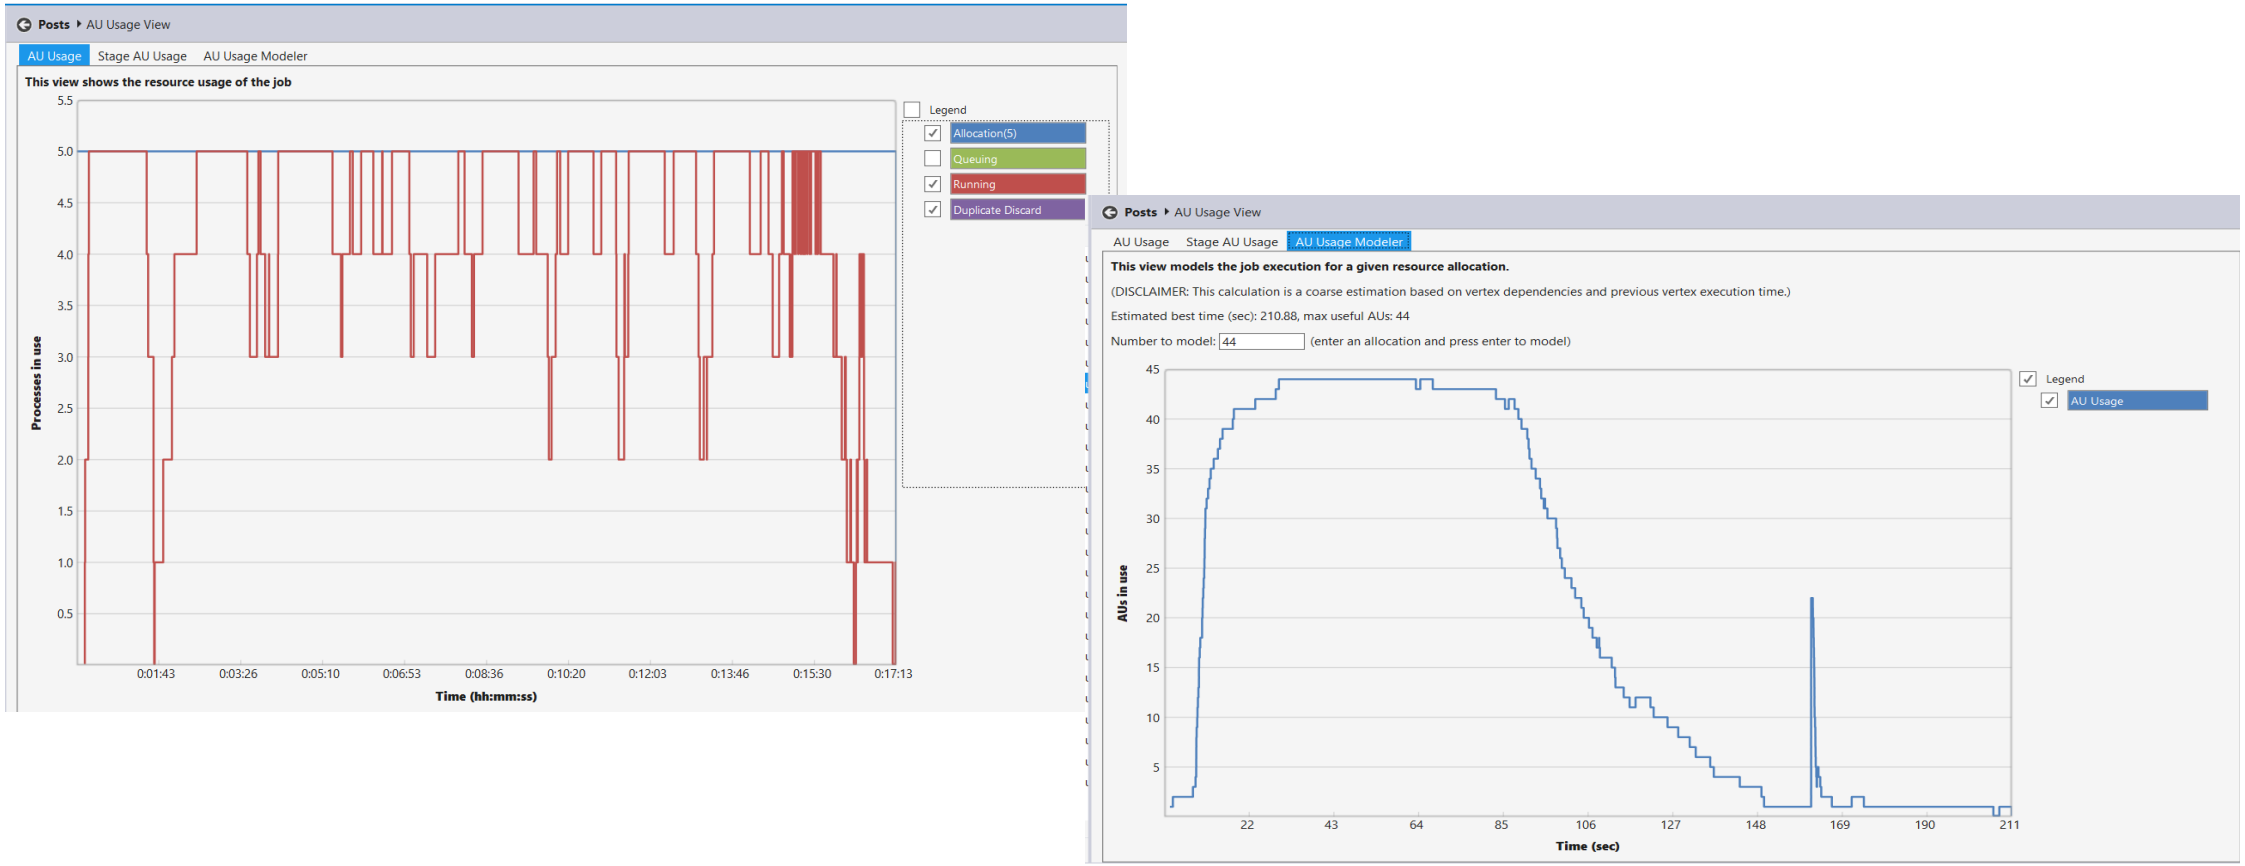

Azure Data Lake Analytics

AZURE DEMO

Azure Data Lake

Q&A